

Effective and Efficient Observability with OpenTelemetry

Daniel Gomez Blanco

Principal Engineer

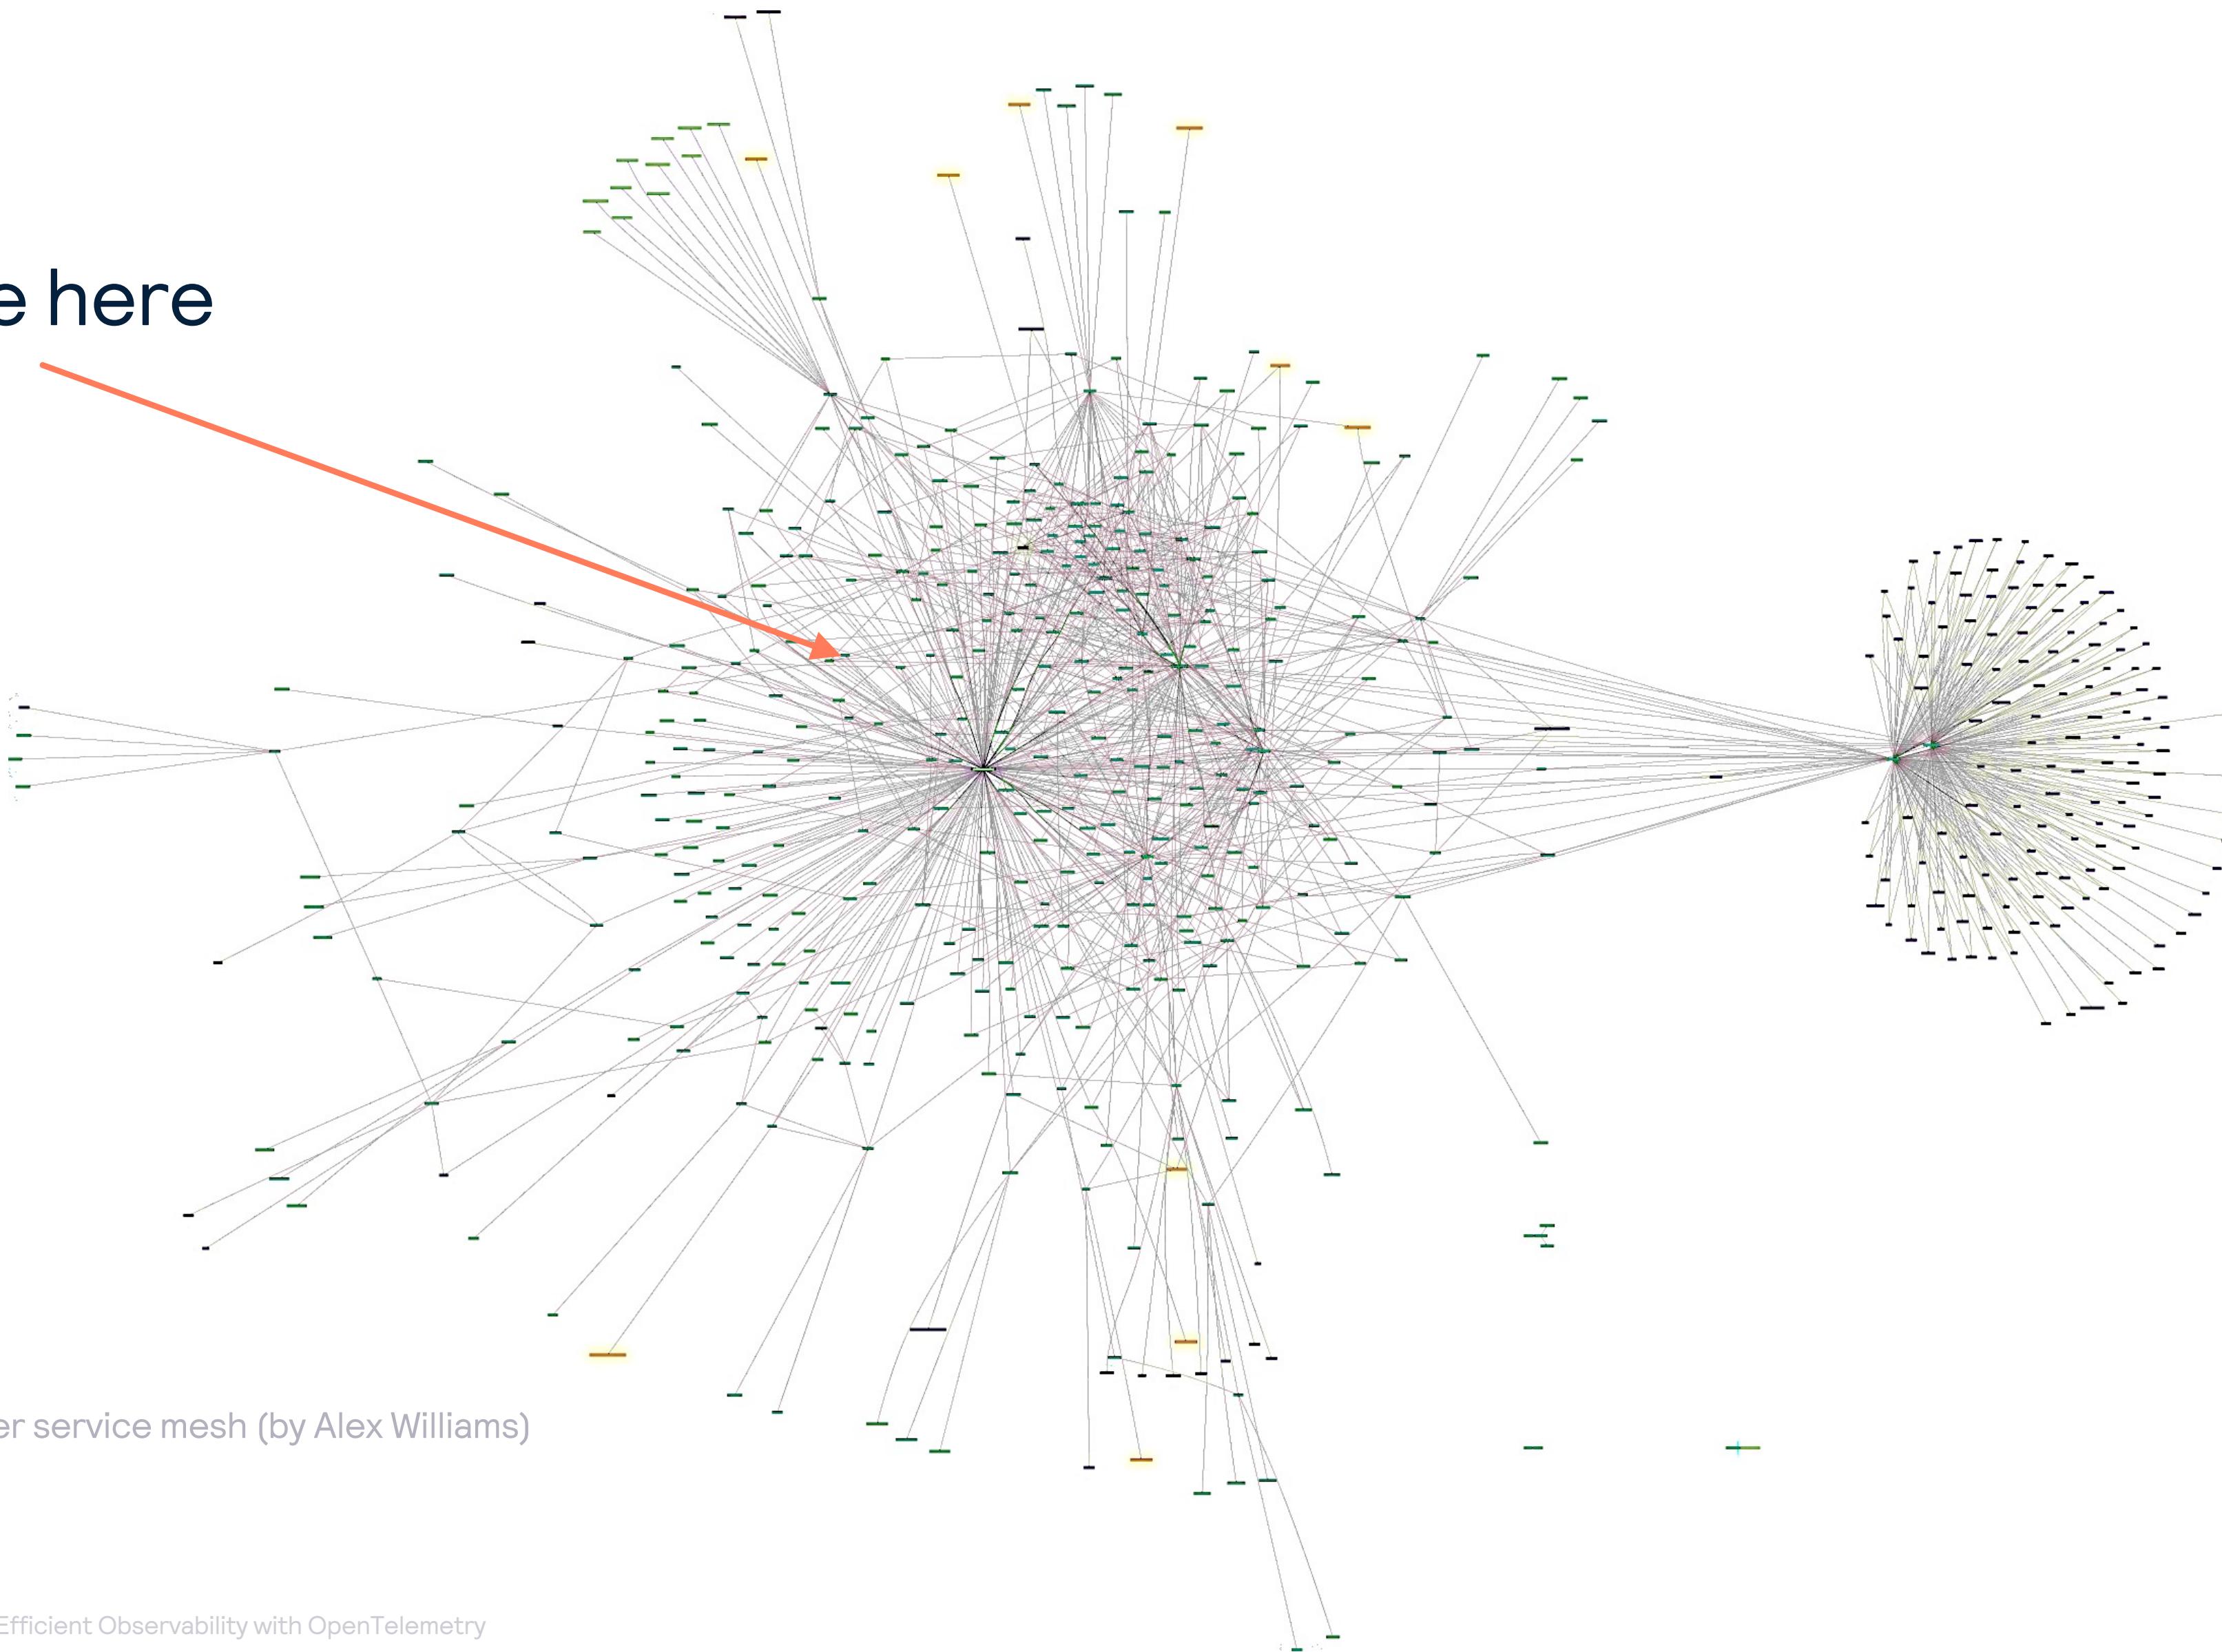
When our systems
change, how
do we know
what changed?

Easy when we are the only ones changing it...

```
if my_condition:  
    print("I'm here")  
    # Do some work  
else:  
    print("I'm there")  
    # Do some other work
```

... not easy in a real distributed system

You're here



Source: service dependencies in Skyscanner service mesh (by Alex Williams)



A bit about me

- Joined Skyscanner in 2018 to work on performance and resource optimisation
- Principal Engineer leading observability strategy since 2020
- 12 years as platform engineer in organisations from 5 to 2,500 employees
- Author of Practical OpenTelemetry: Adopting Open Observability Standards Across Your Organization (Apress, 2023)

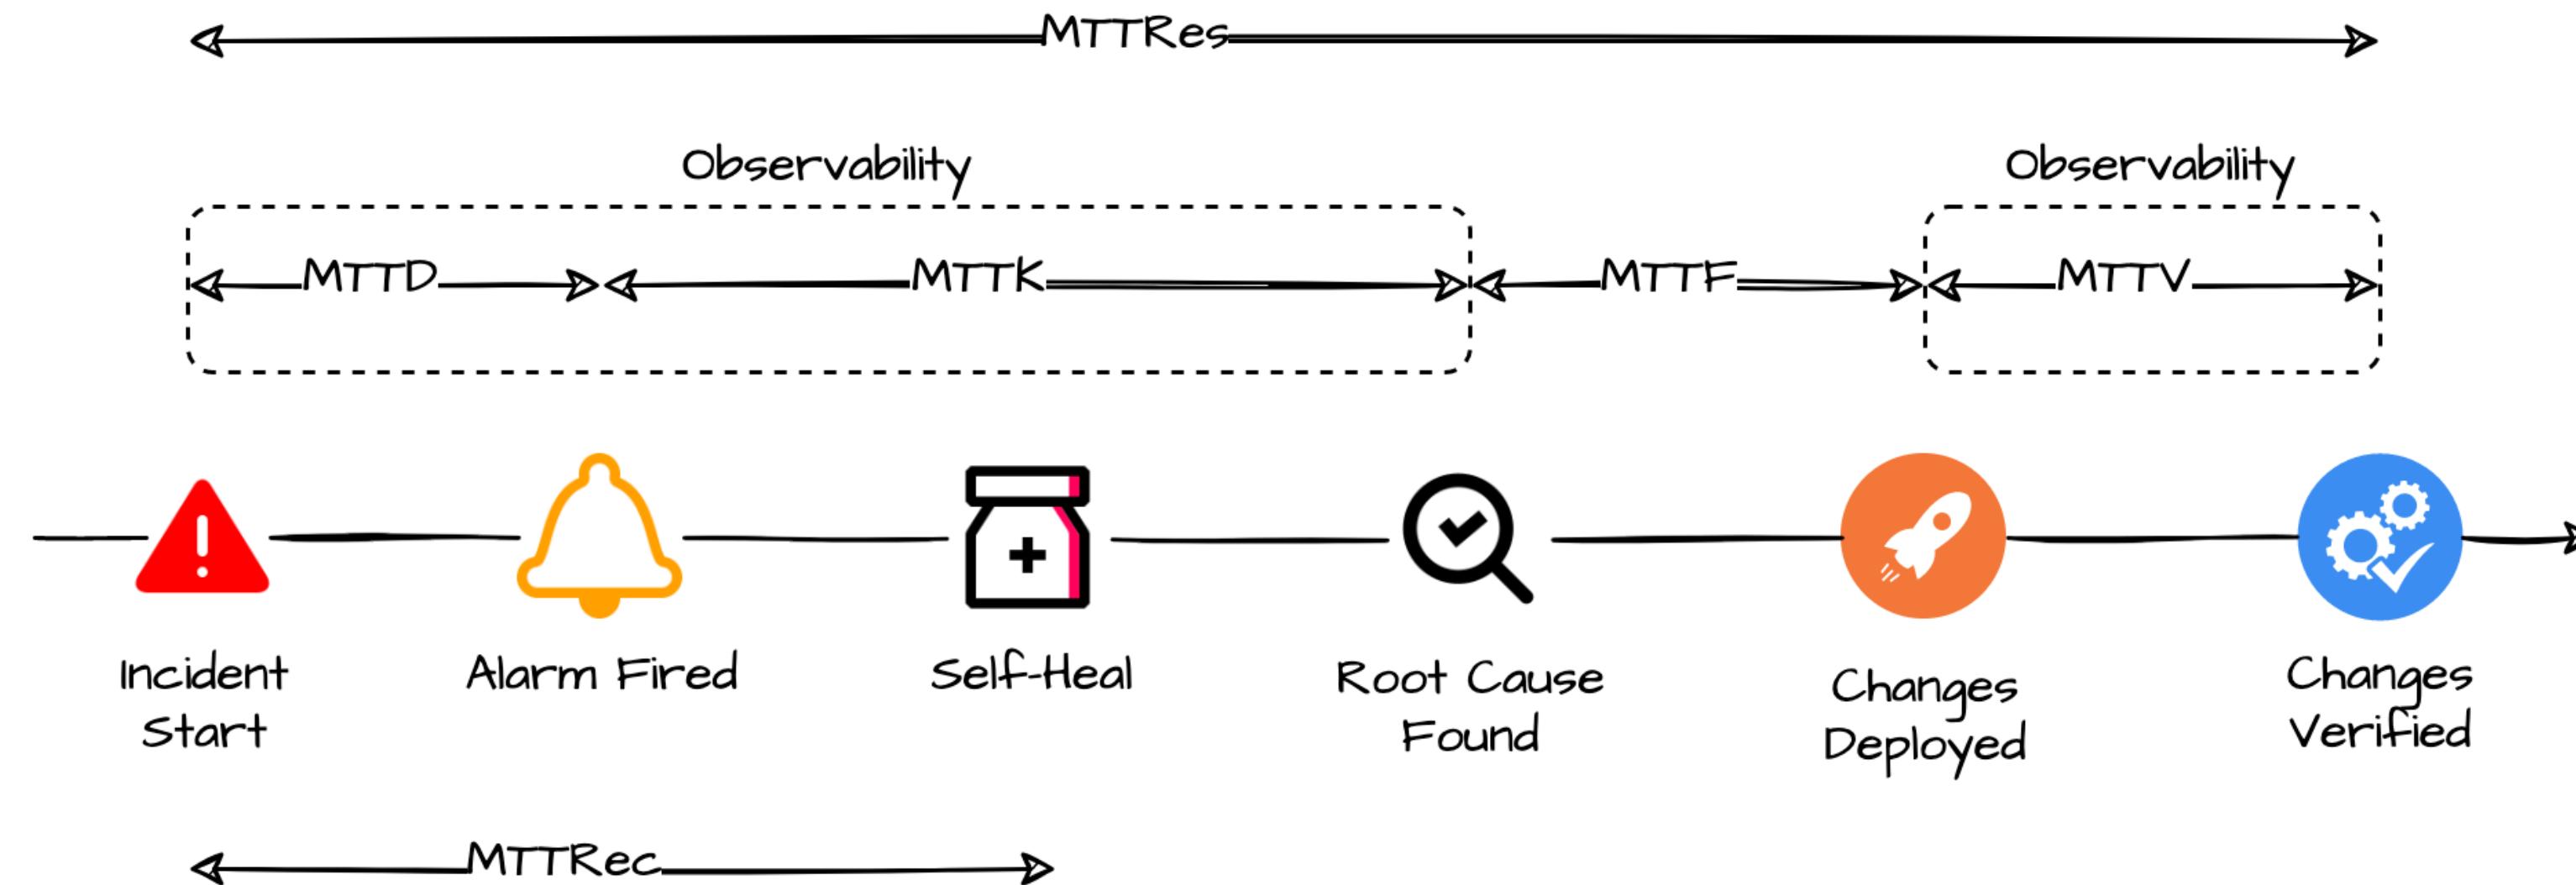
Agenda

- 1. Why observability matters**
- 2. How open standards help observability**
- 3. Rolling out OpenTelemetry**
- 4. Adopting observability in practice**

Agenda

- 1. Why observability matters**
2. How open standards help observability
3. Rolling out OpenTelemetry
4. Adopting observability in practice

Observability within incident response



Source: Practical OpenTelemetry: Adopting Open Observability Standards Across Your Organization

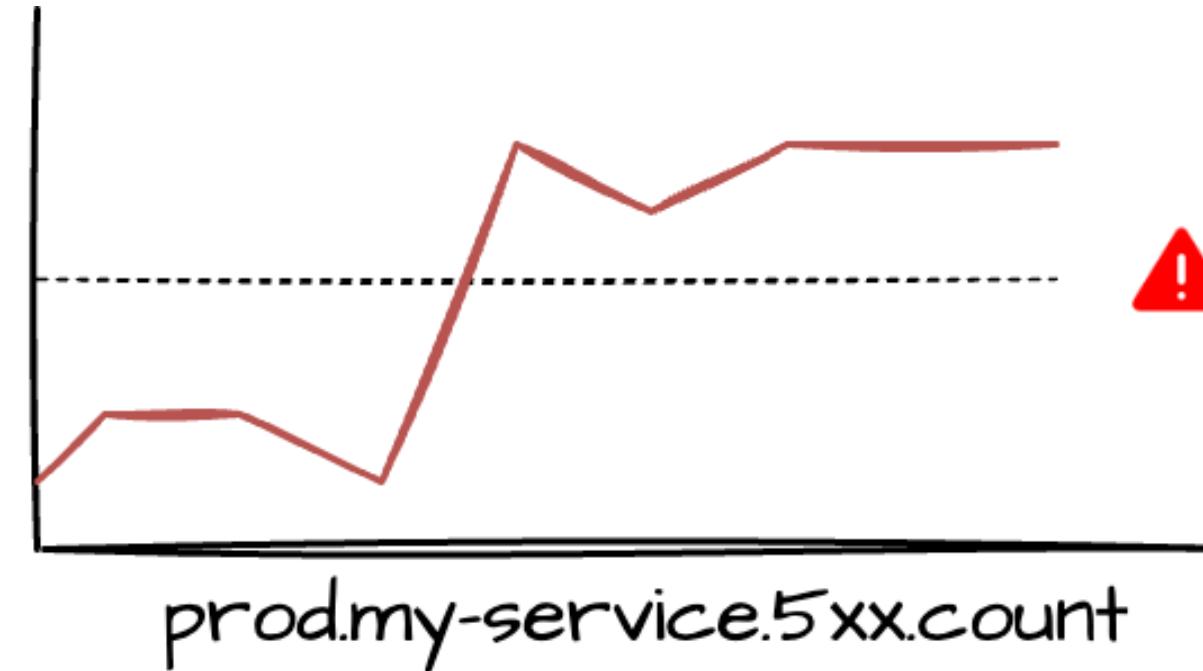
Is my system behaving as expected?

Reducing Mean-Time-to-Detect and Mean-Time-To-Verify

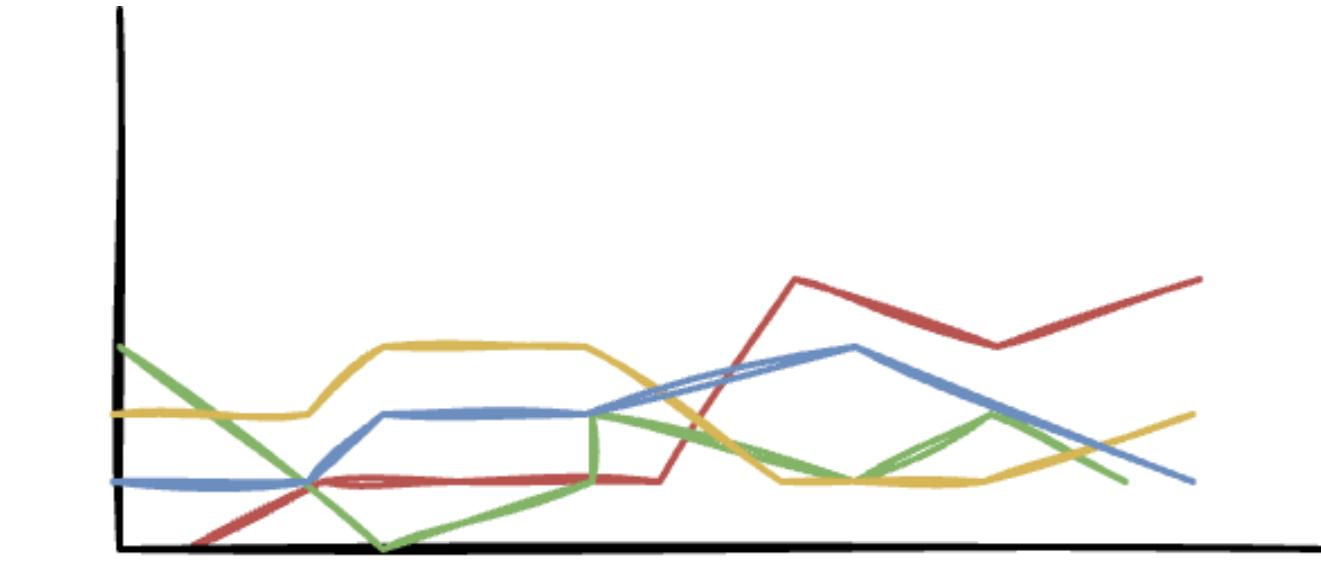
Why is my system not behaving as expected?

Reducing Mean-Time-to-Know

Debugging that relies on past experience



Follow runbook to dashboard



prod.my-service.memory.usage.gauge

Search for known ERROR cases
↓

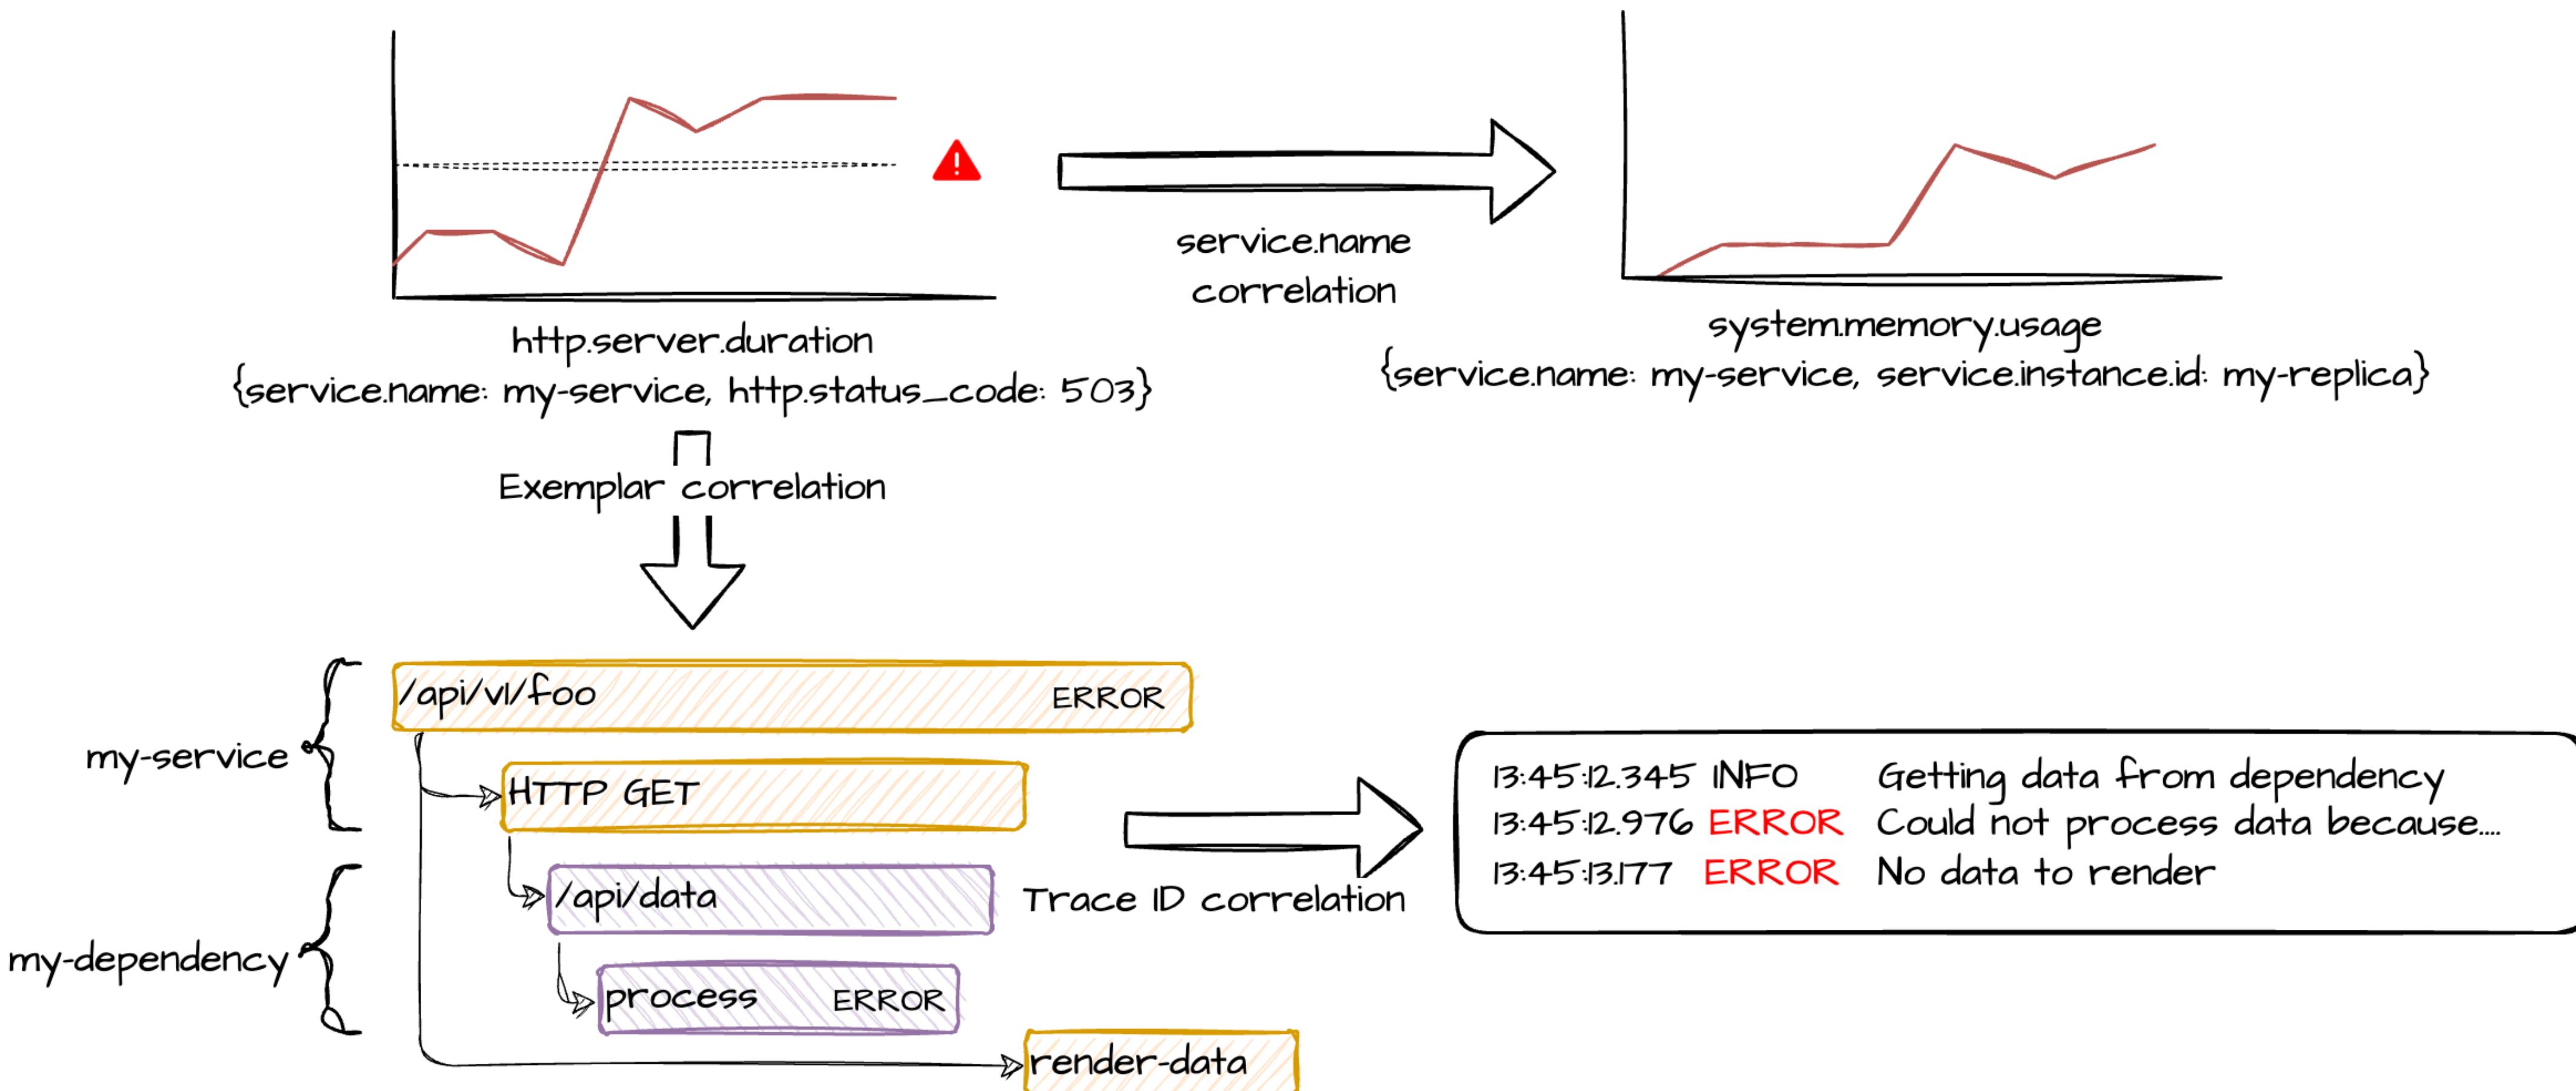
13:45:13.177 **ERROR** No data to render

Call another team??



Now what?

Debugging that relies on context



Effective observability means...

High granularity

Detailed telemetry data corresponding to individual operations within system transactions

Rich context

Considering multiple telemetry signals and dependencies under one single holistic view of the system

Signal correlation

Linking metrics, traces and logs under one single stream of events

Service correlation

Relating telemetry from different services part of the same common operation

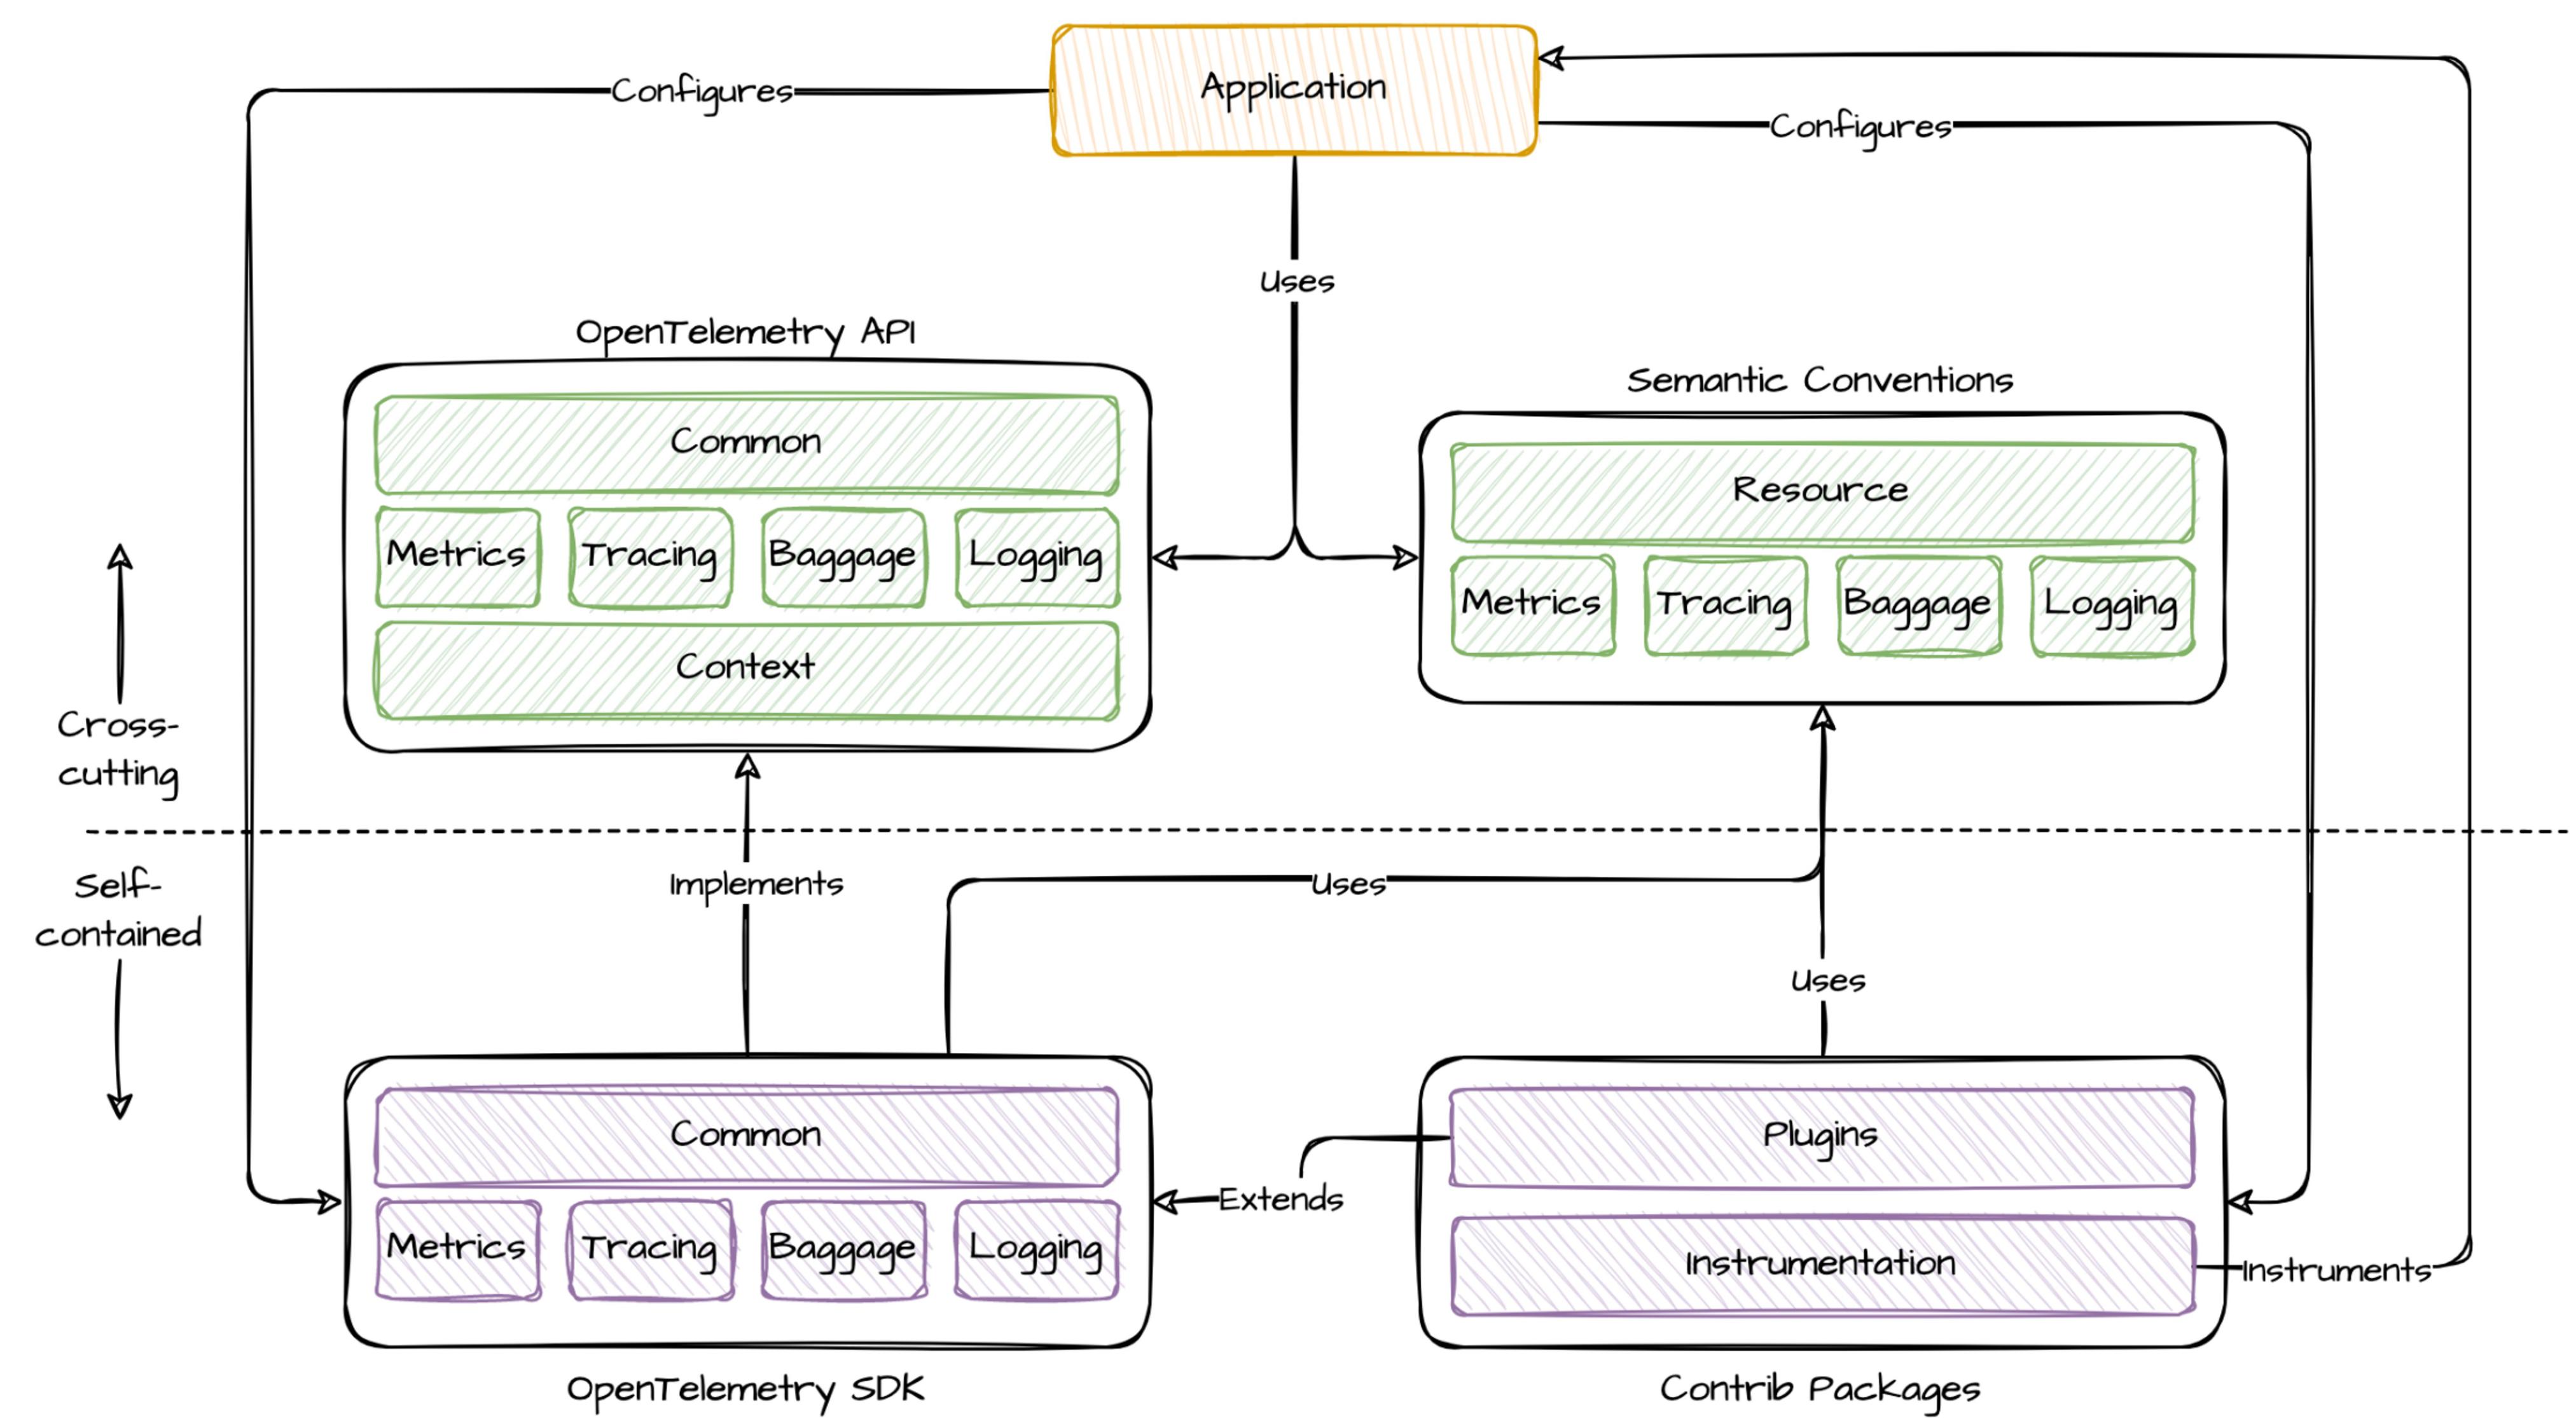
Open standards

Out-of-the-box telemetry instrumented by experts, following open standards across all platforms

Agenda

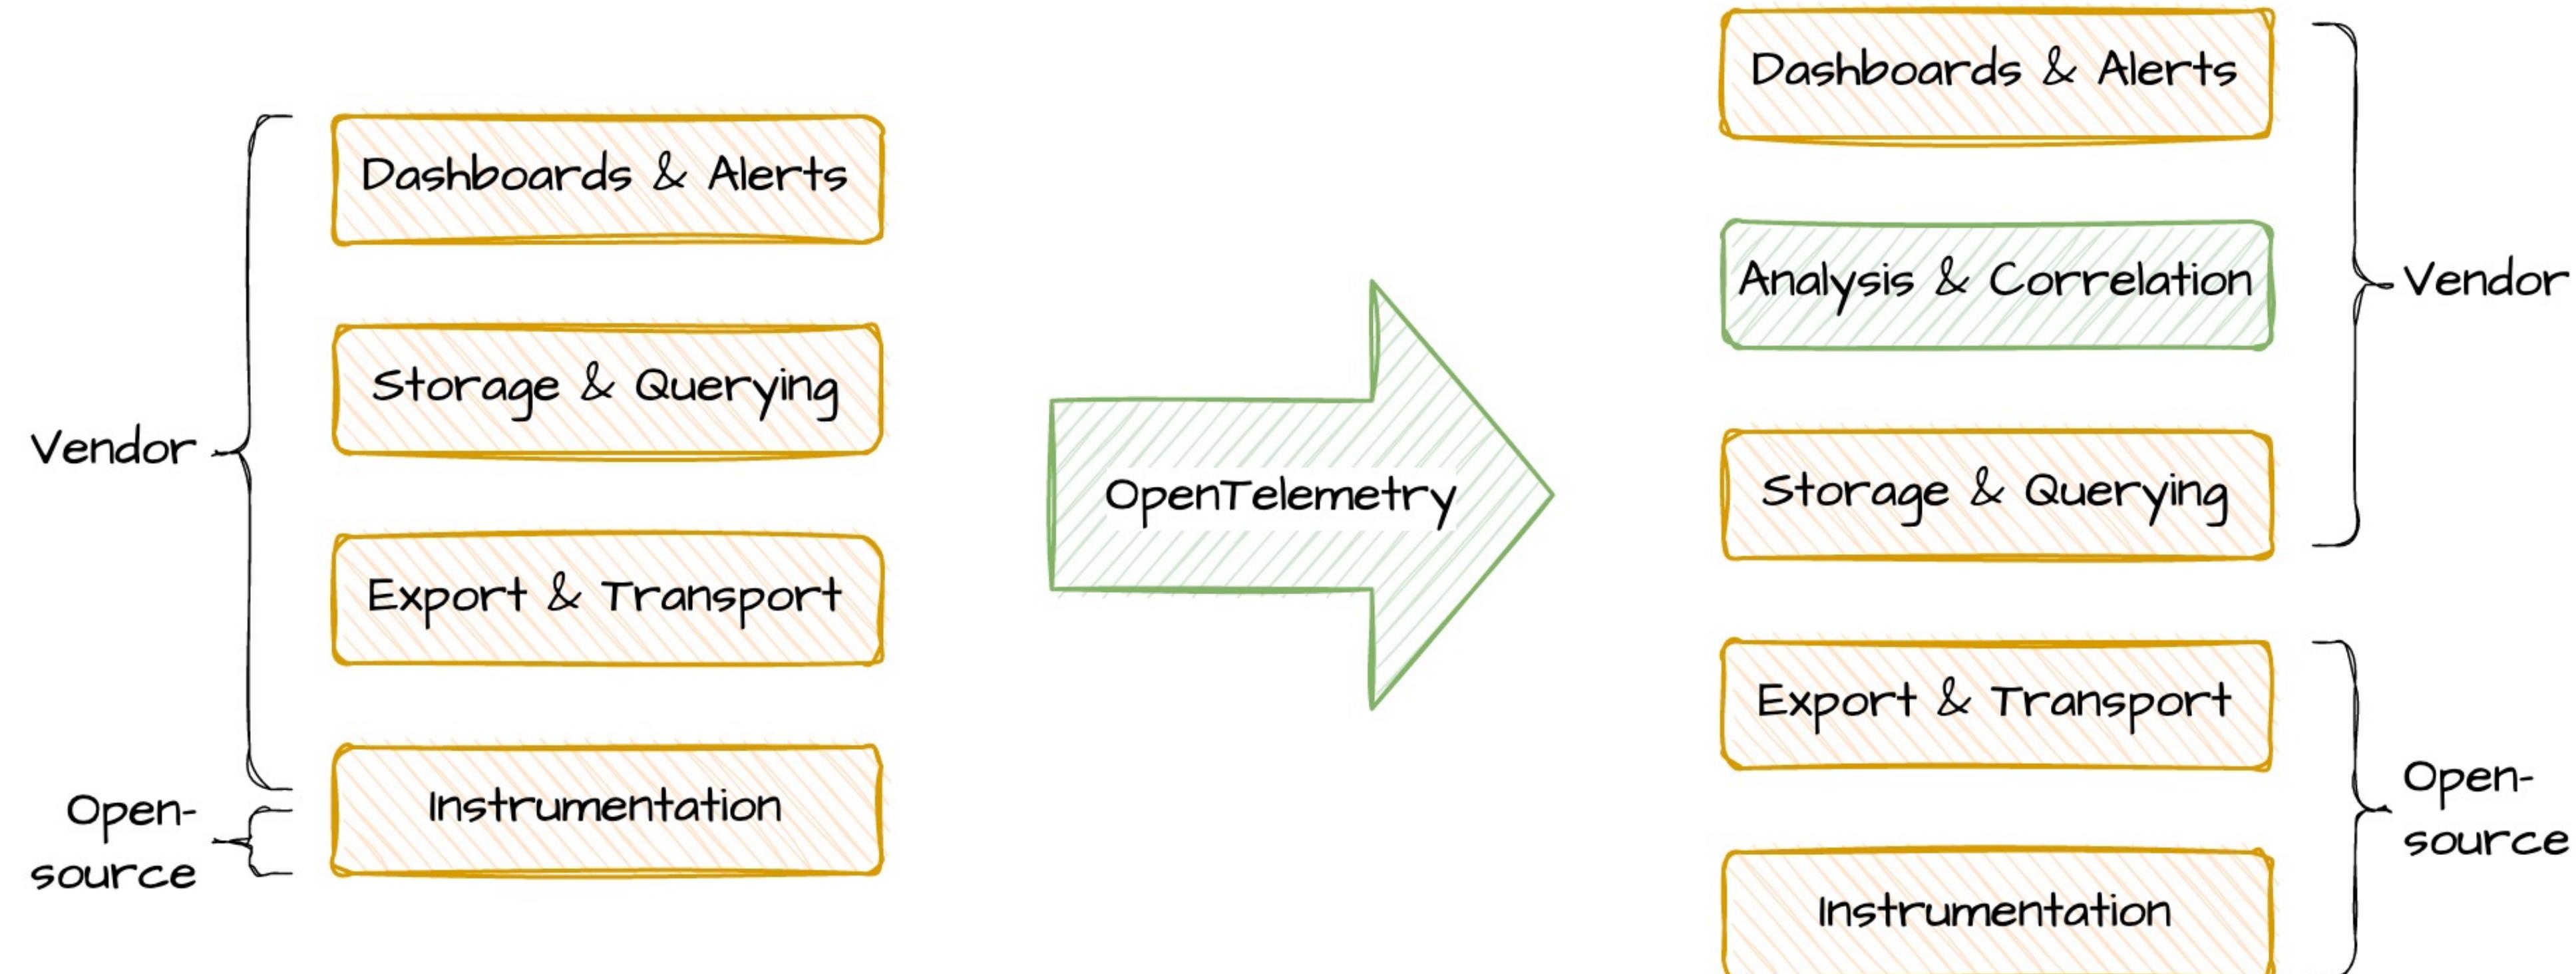
1. Why observability matters
- 2. How open standards help observability**
3. Telemetry signals and their purpose
4. Adopting observability in practice

To enable effective observability by making high-quality, portable telemetry ubiquitous



Practical OpenTelemetry: Adopting Open Observability Standards Across Your Organization

Influence in buy-vs-build decisions

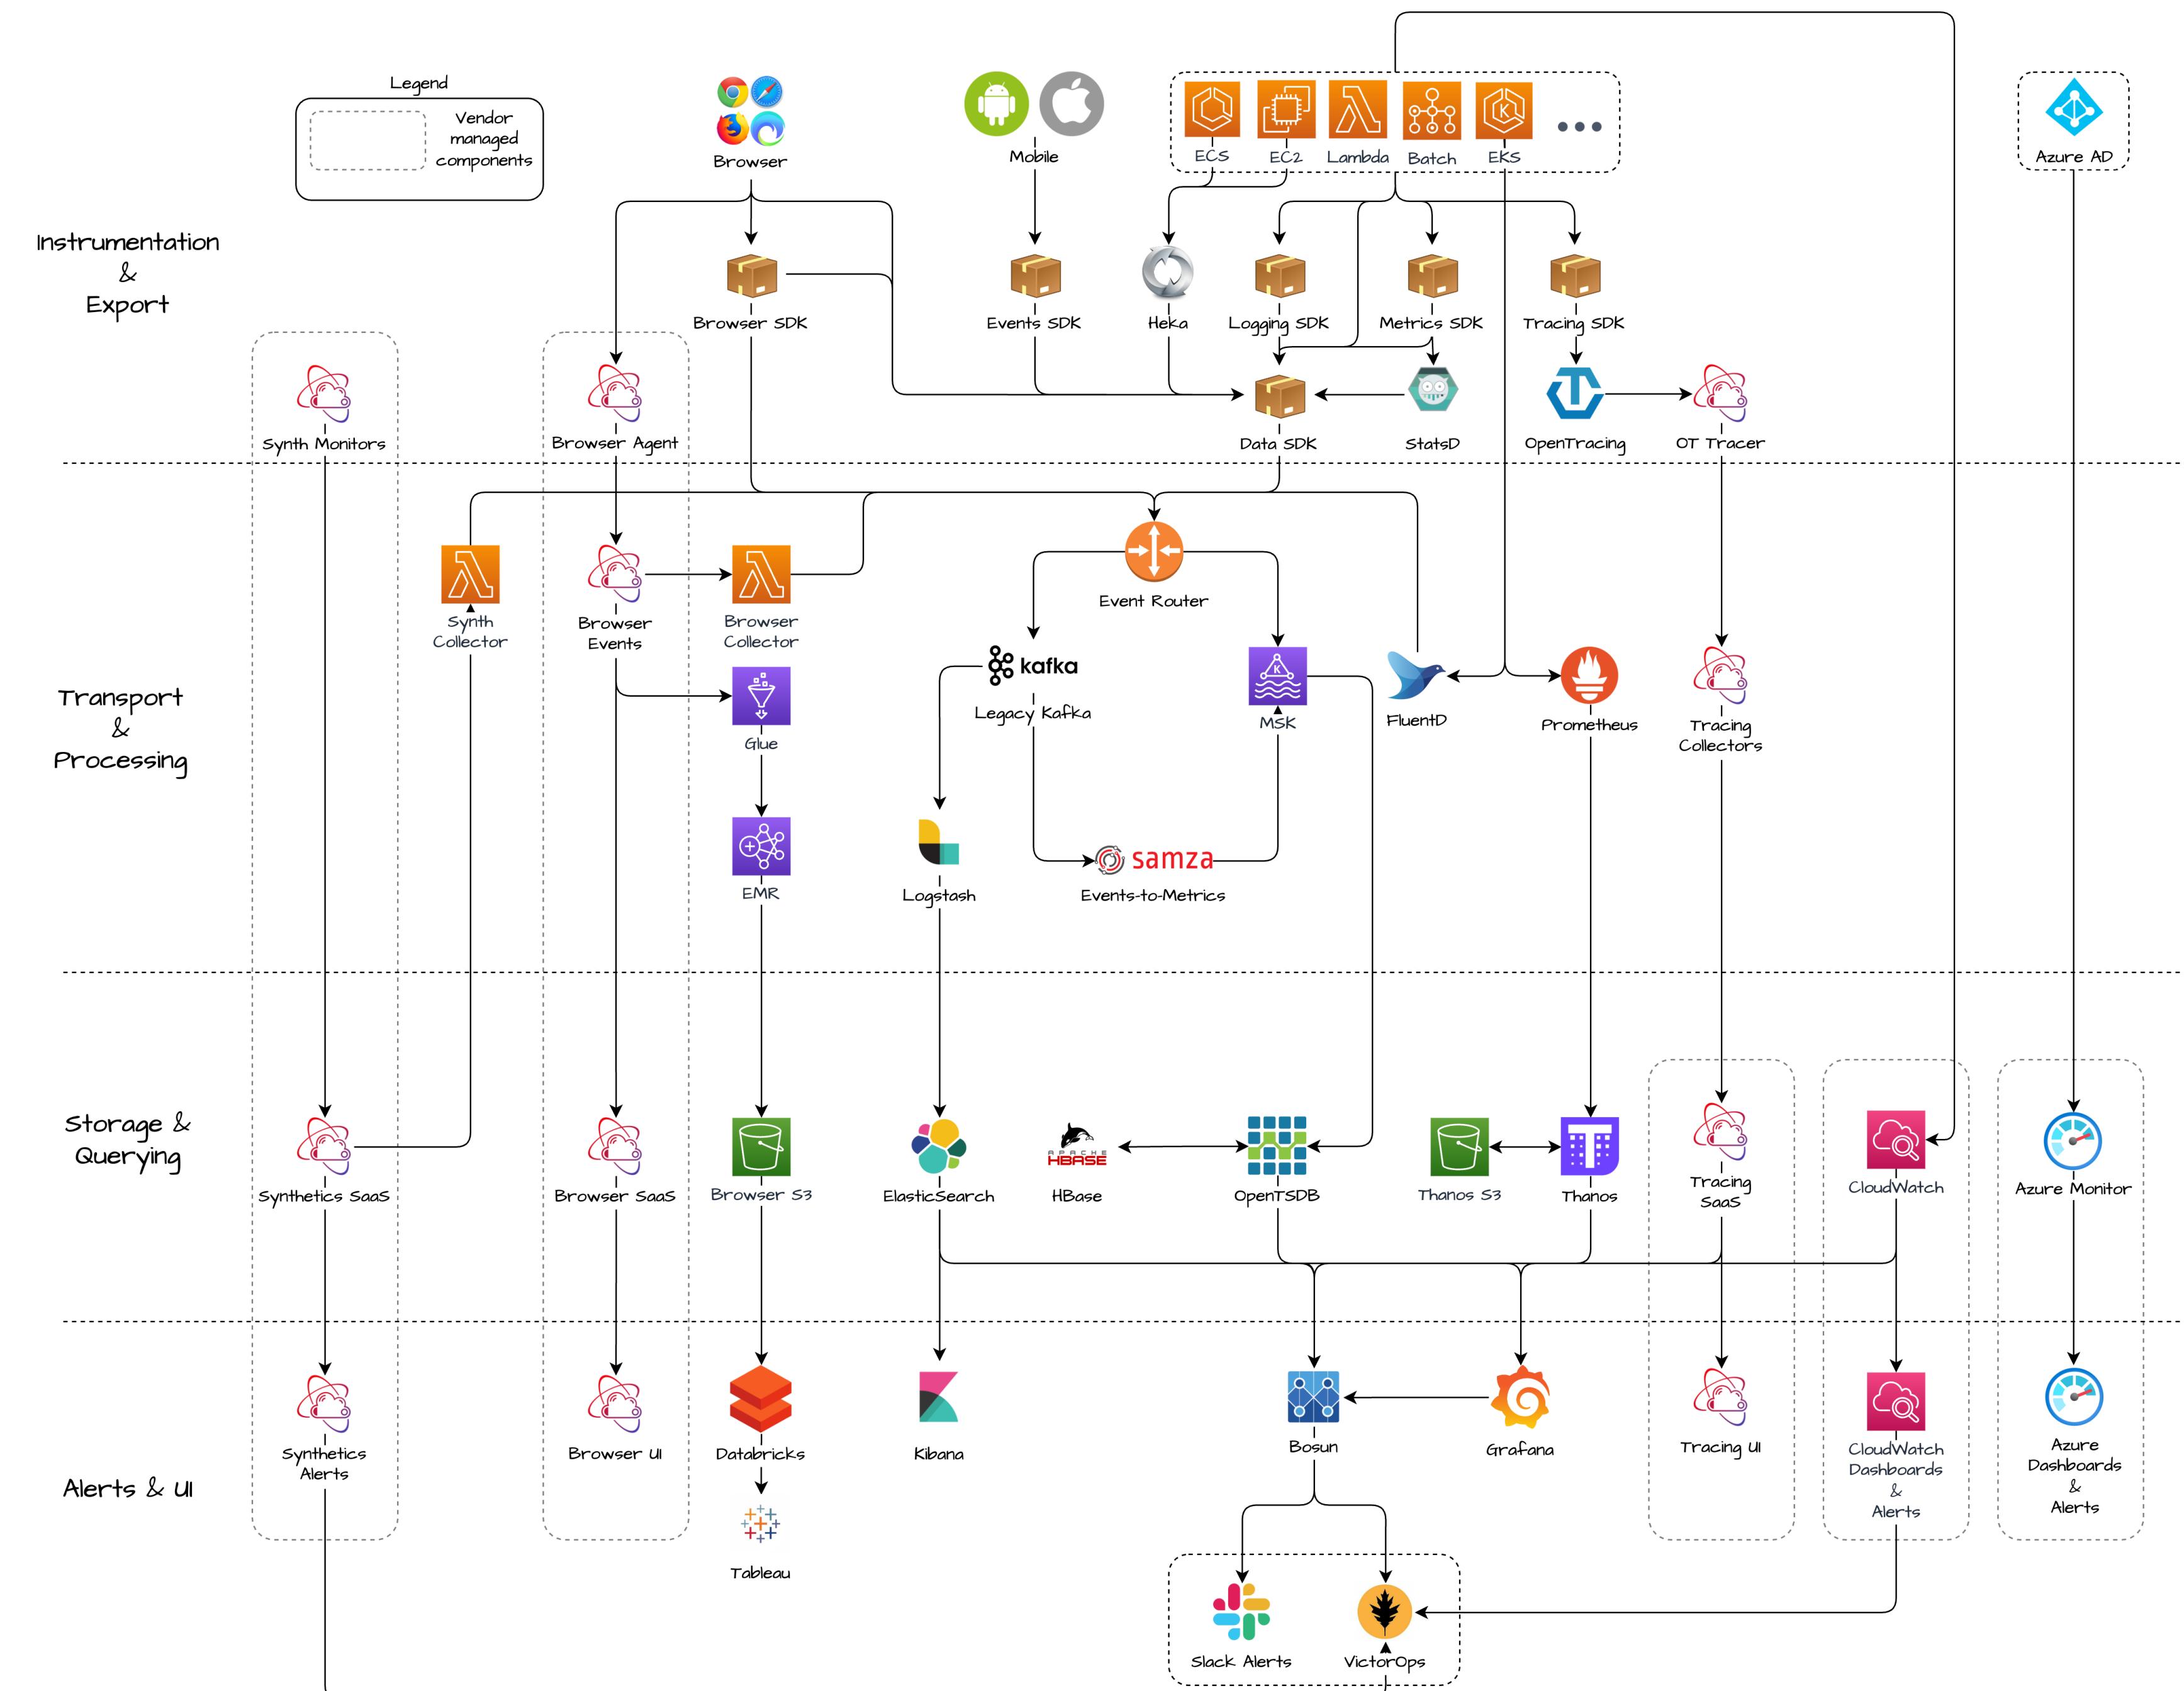


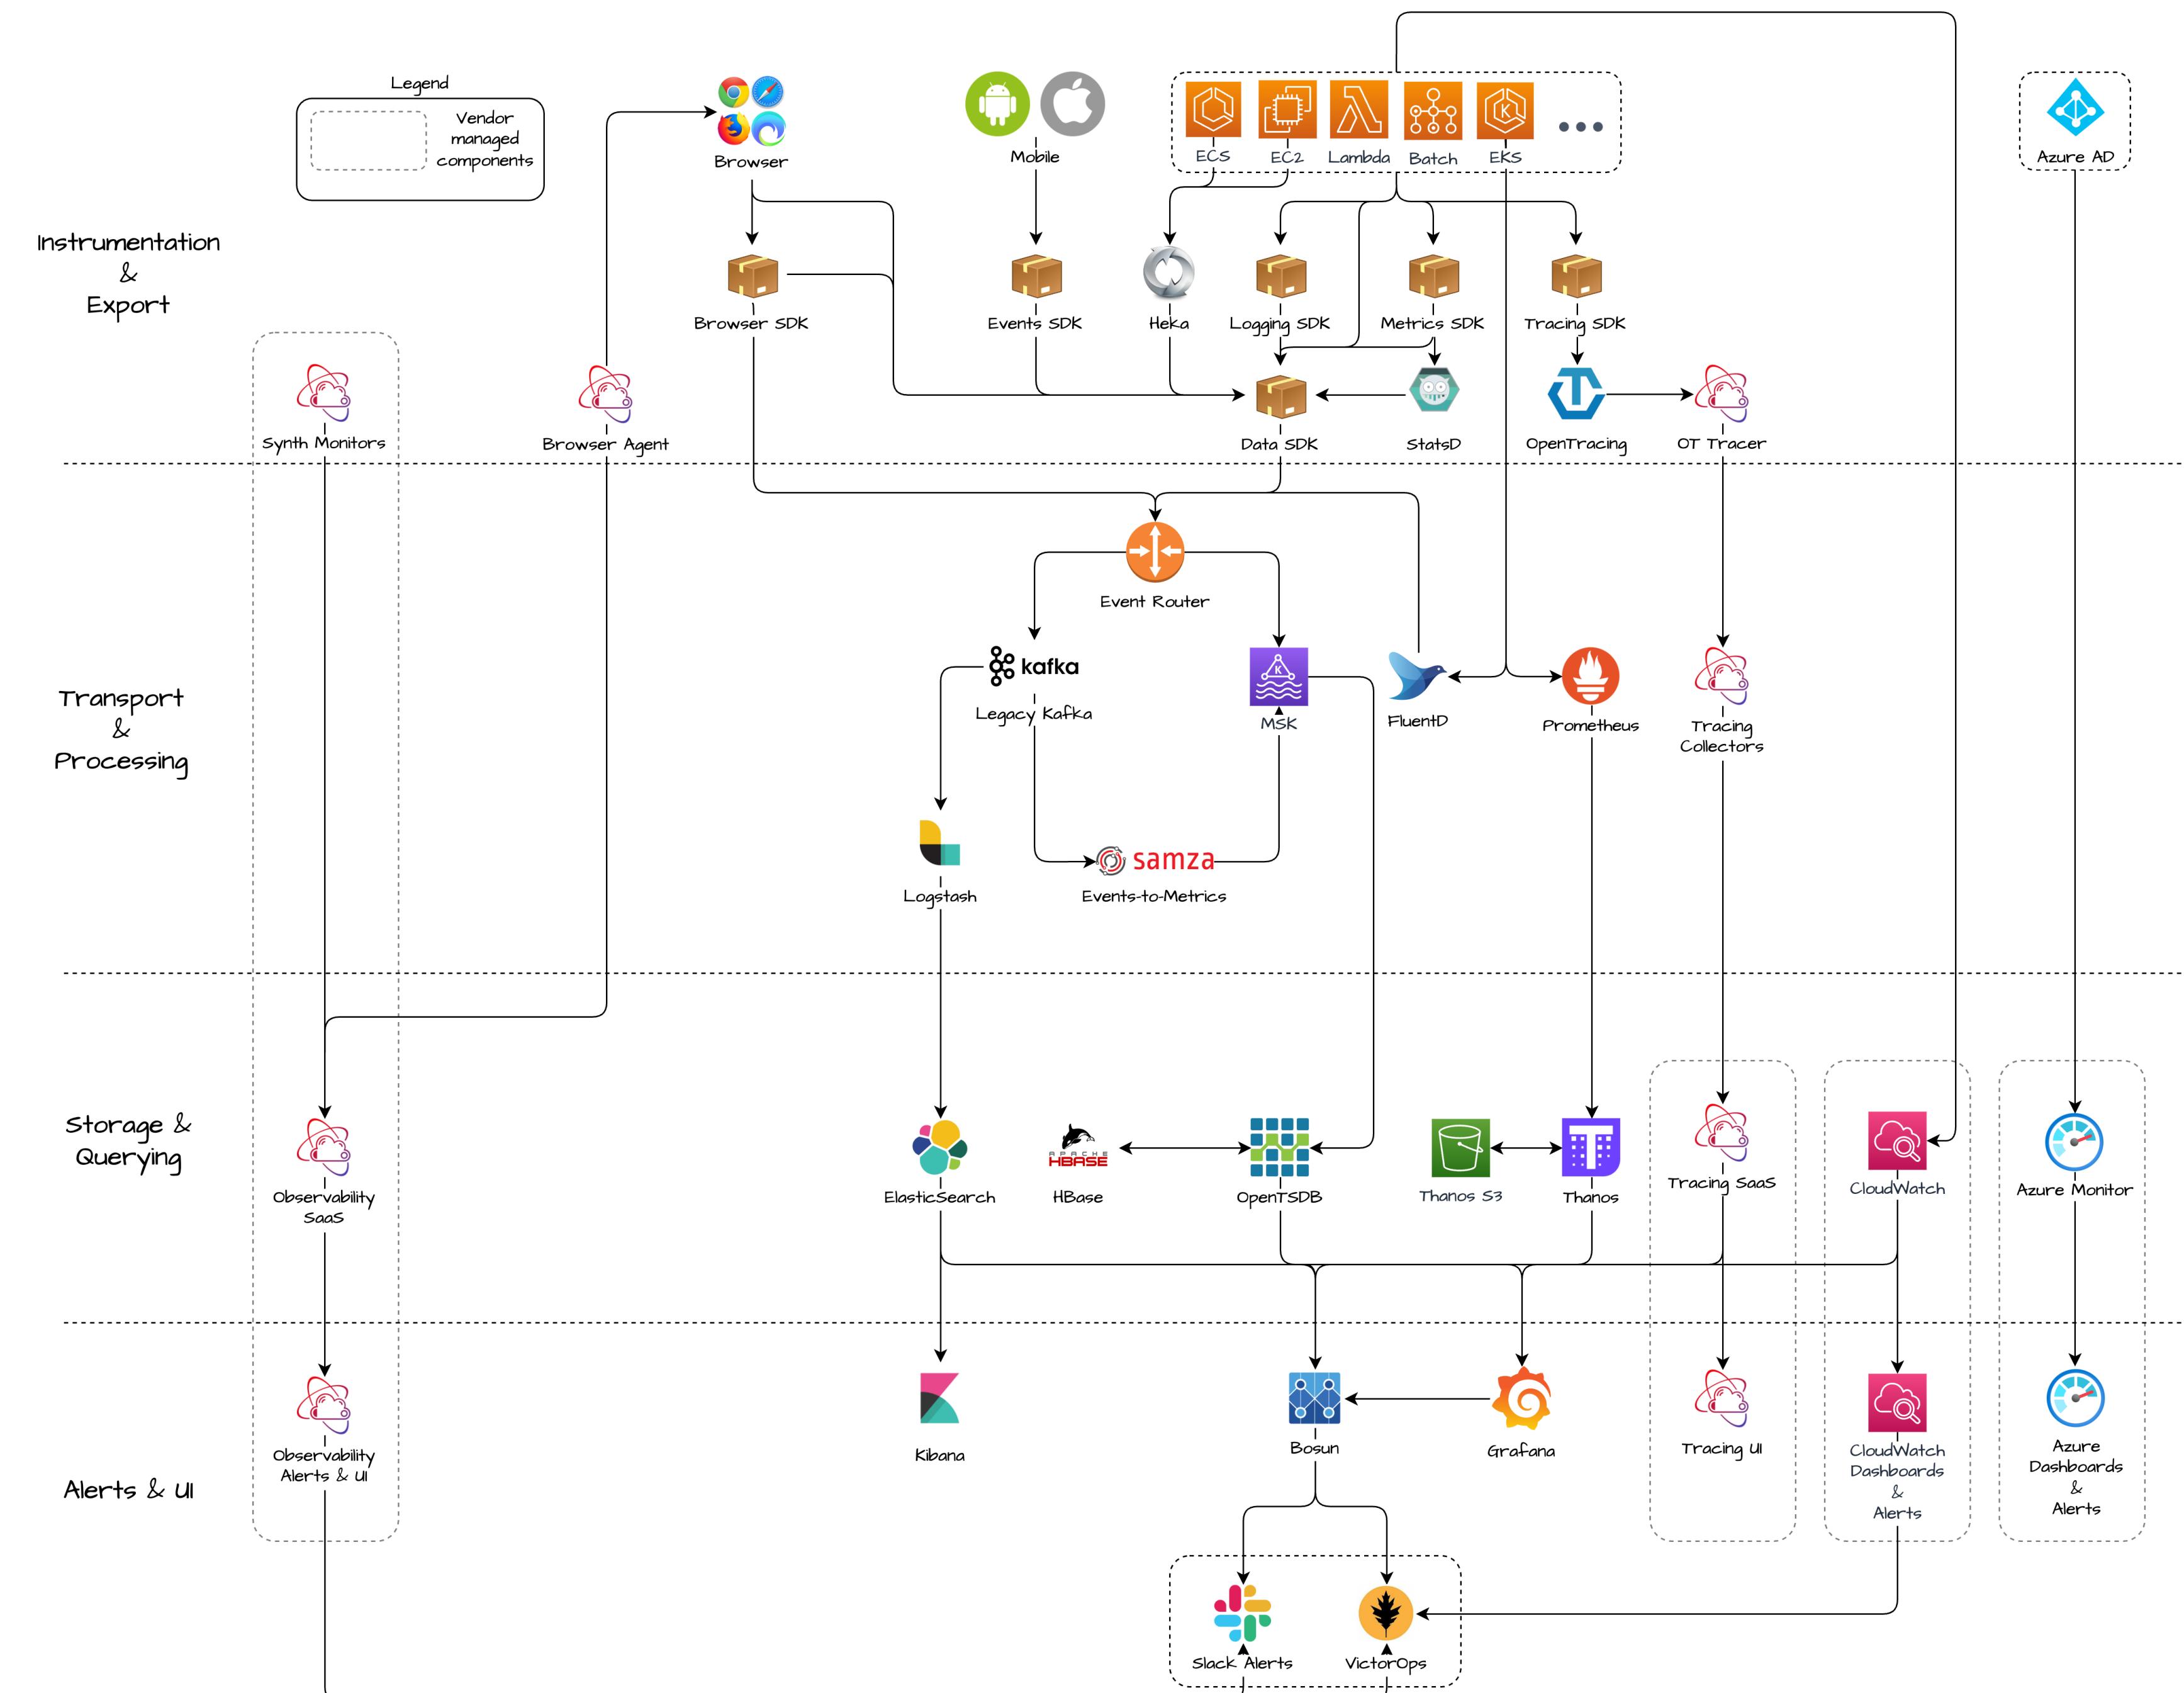
Practical OpenTelemetry: Adopting Open Observability Standards Across Your Organization

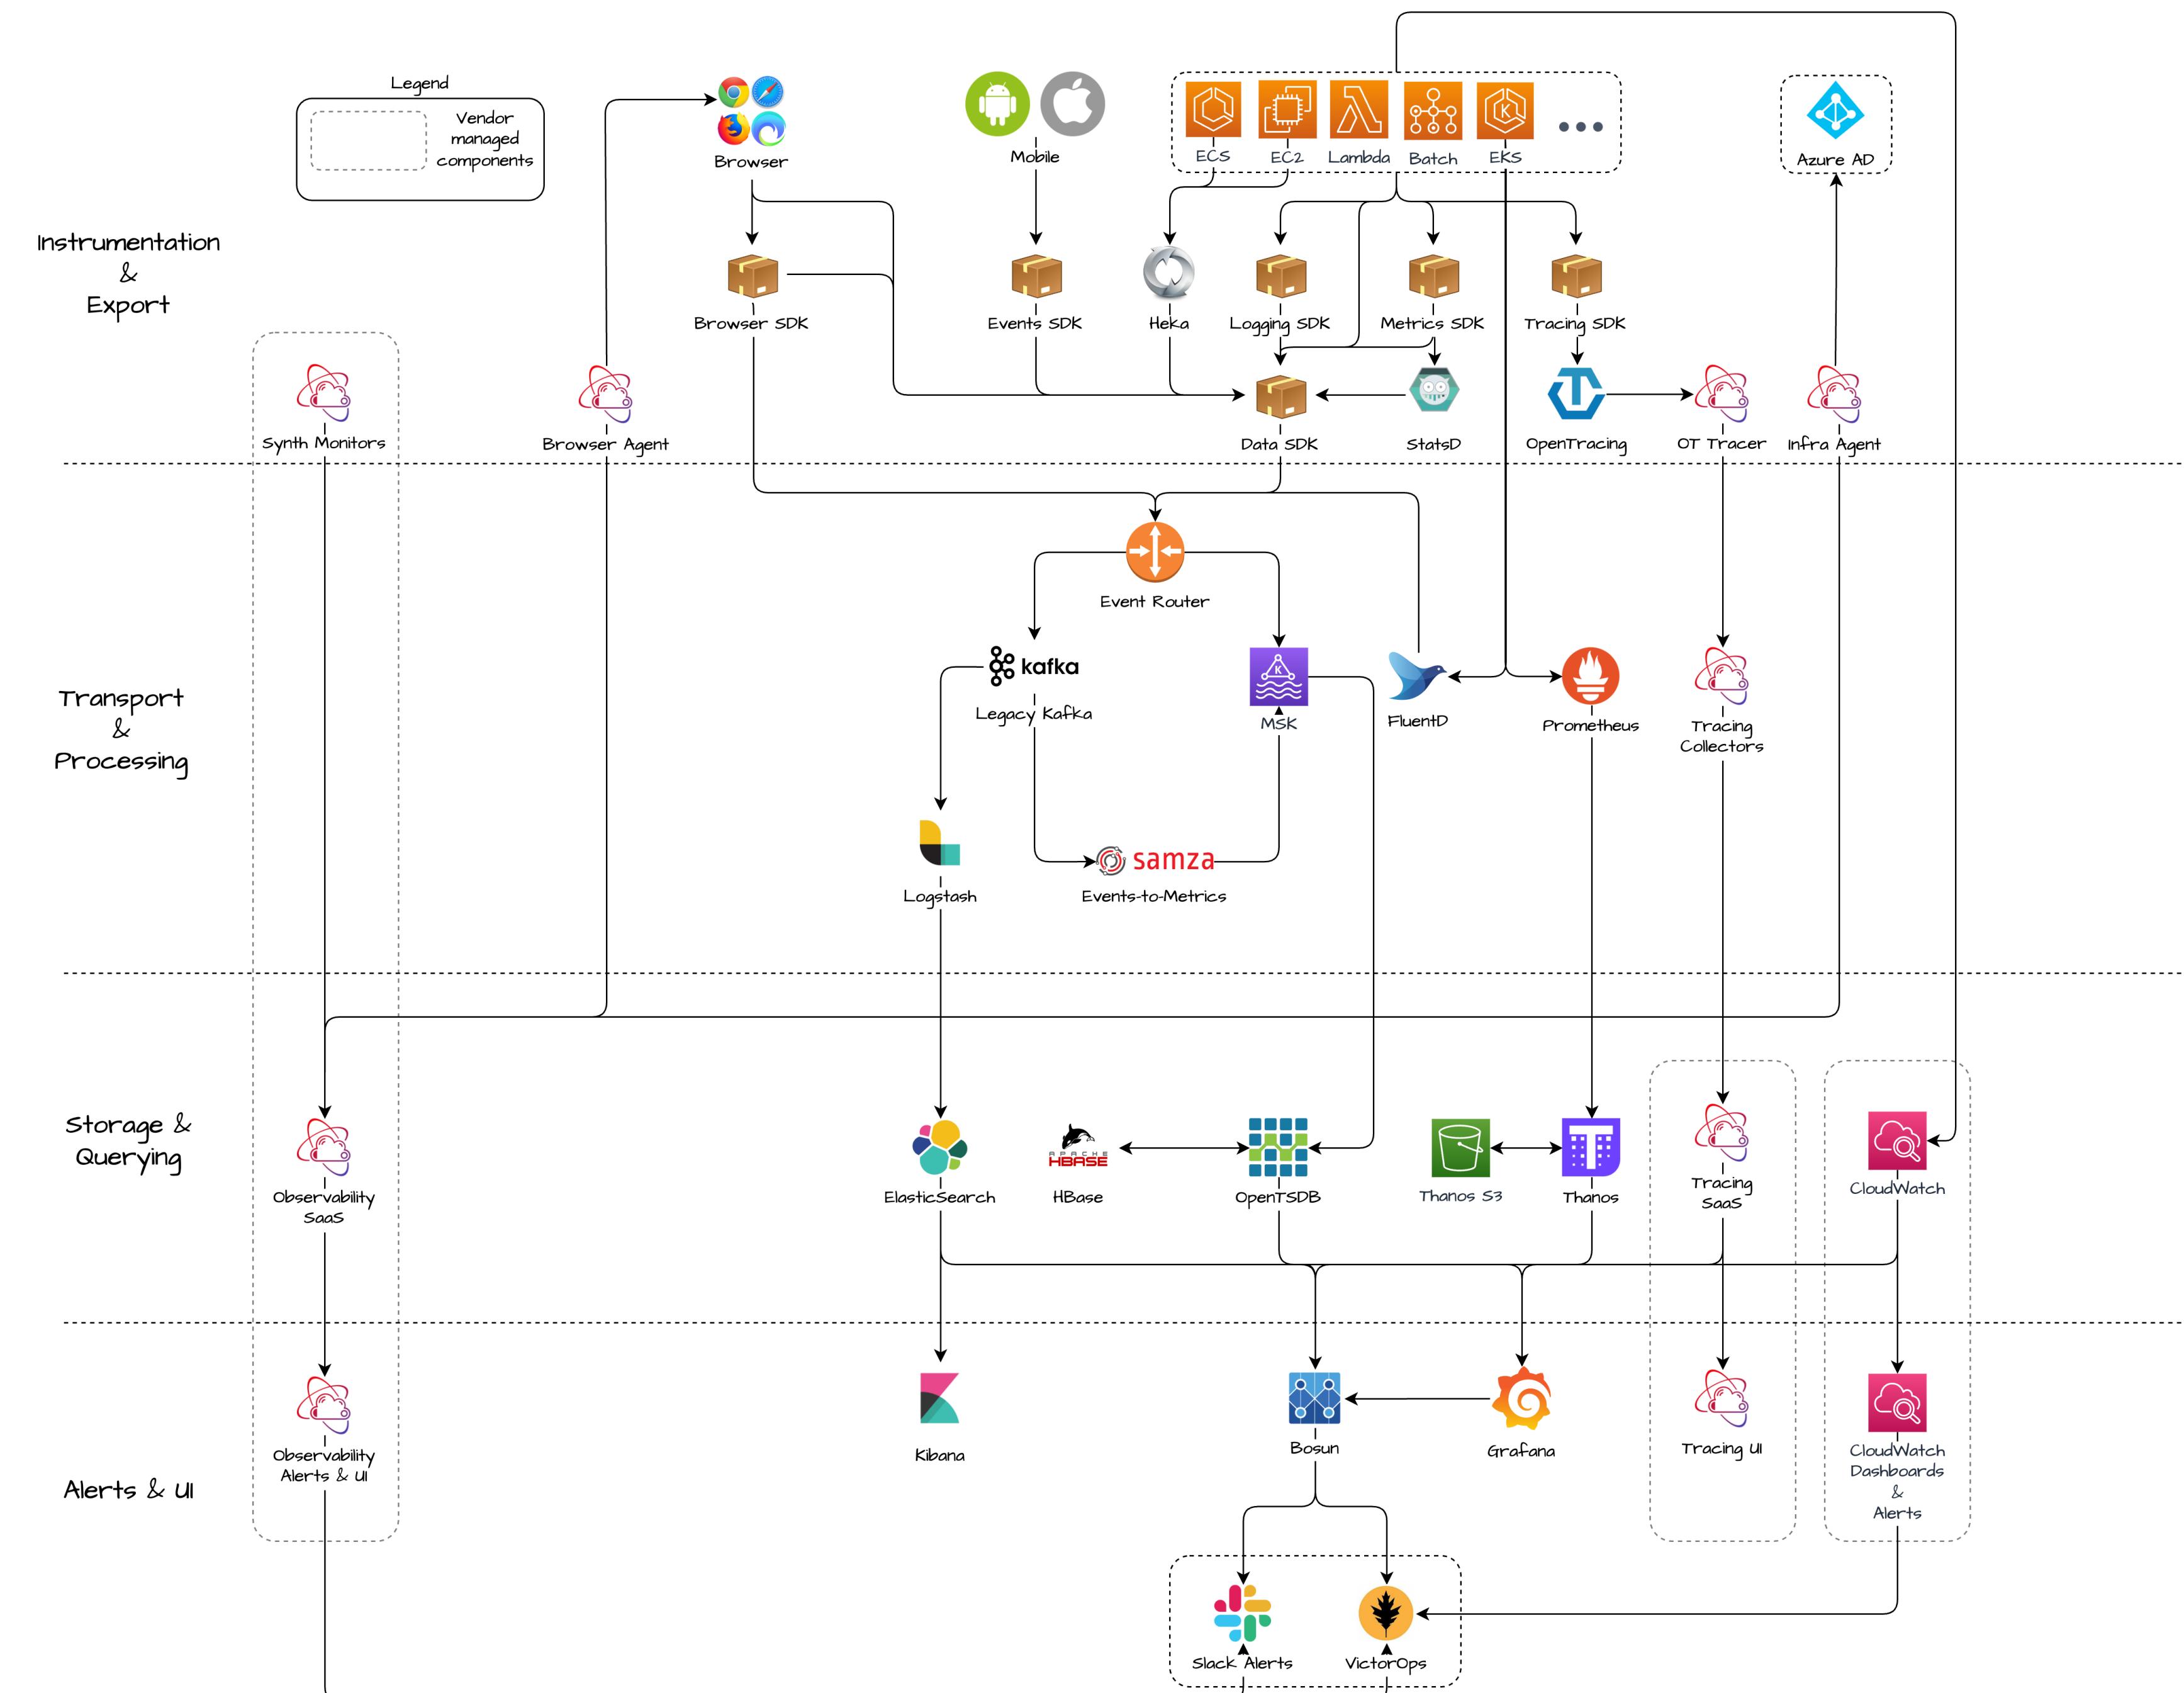
Agenda

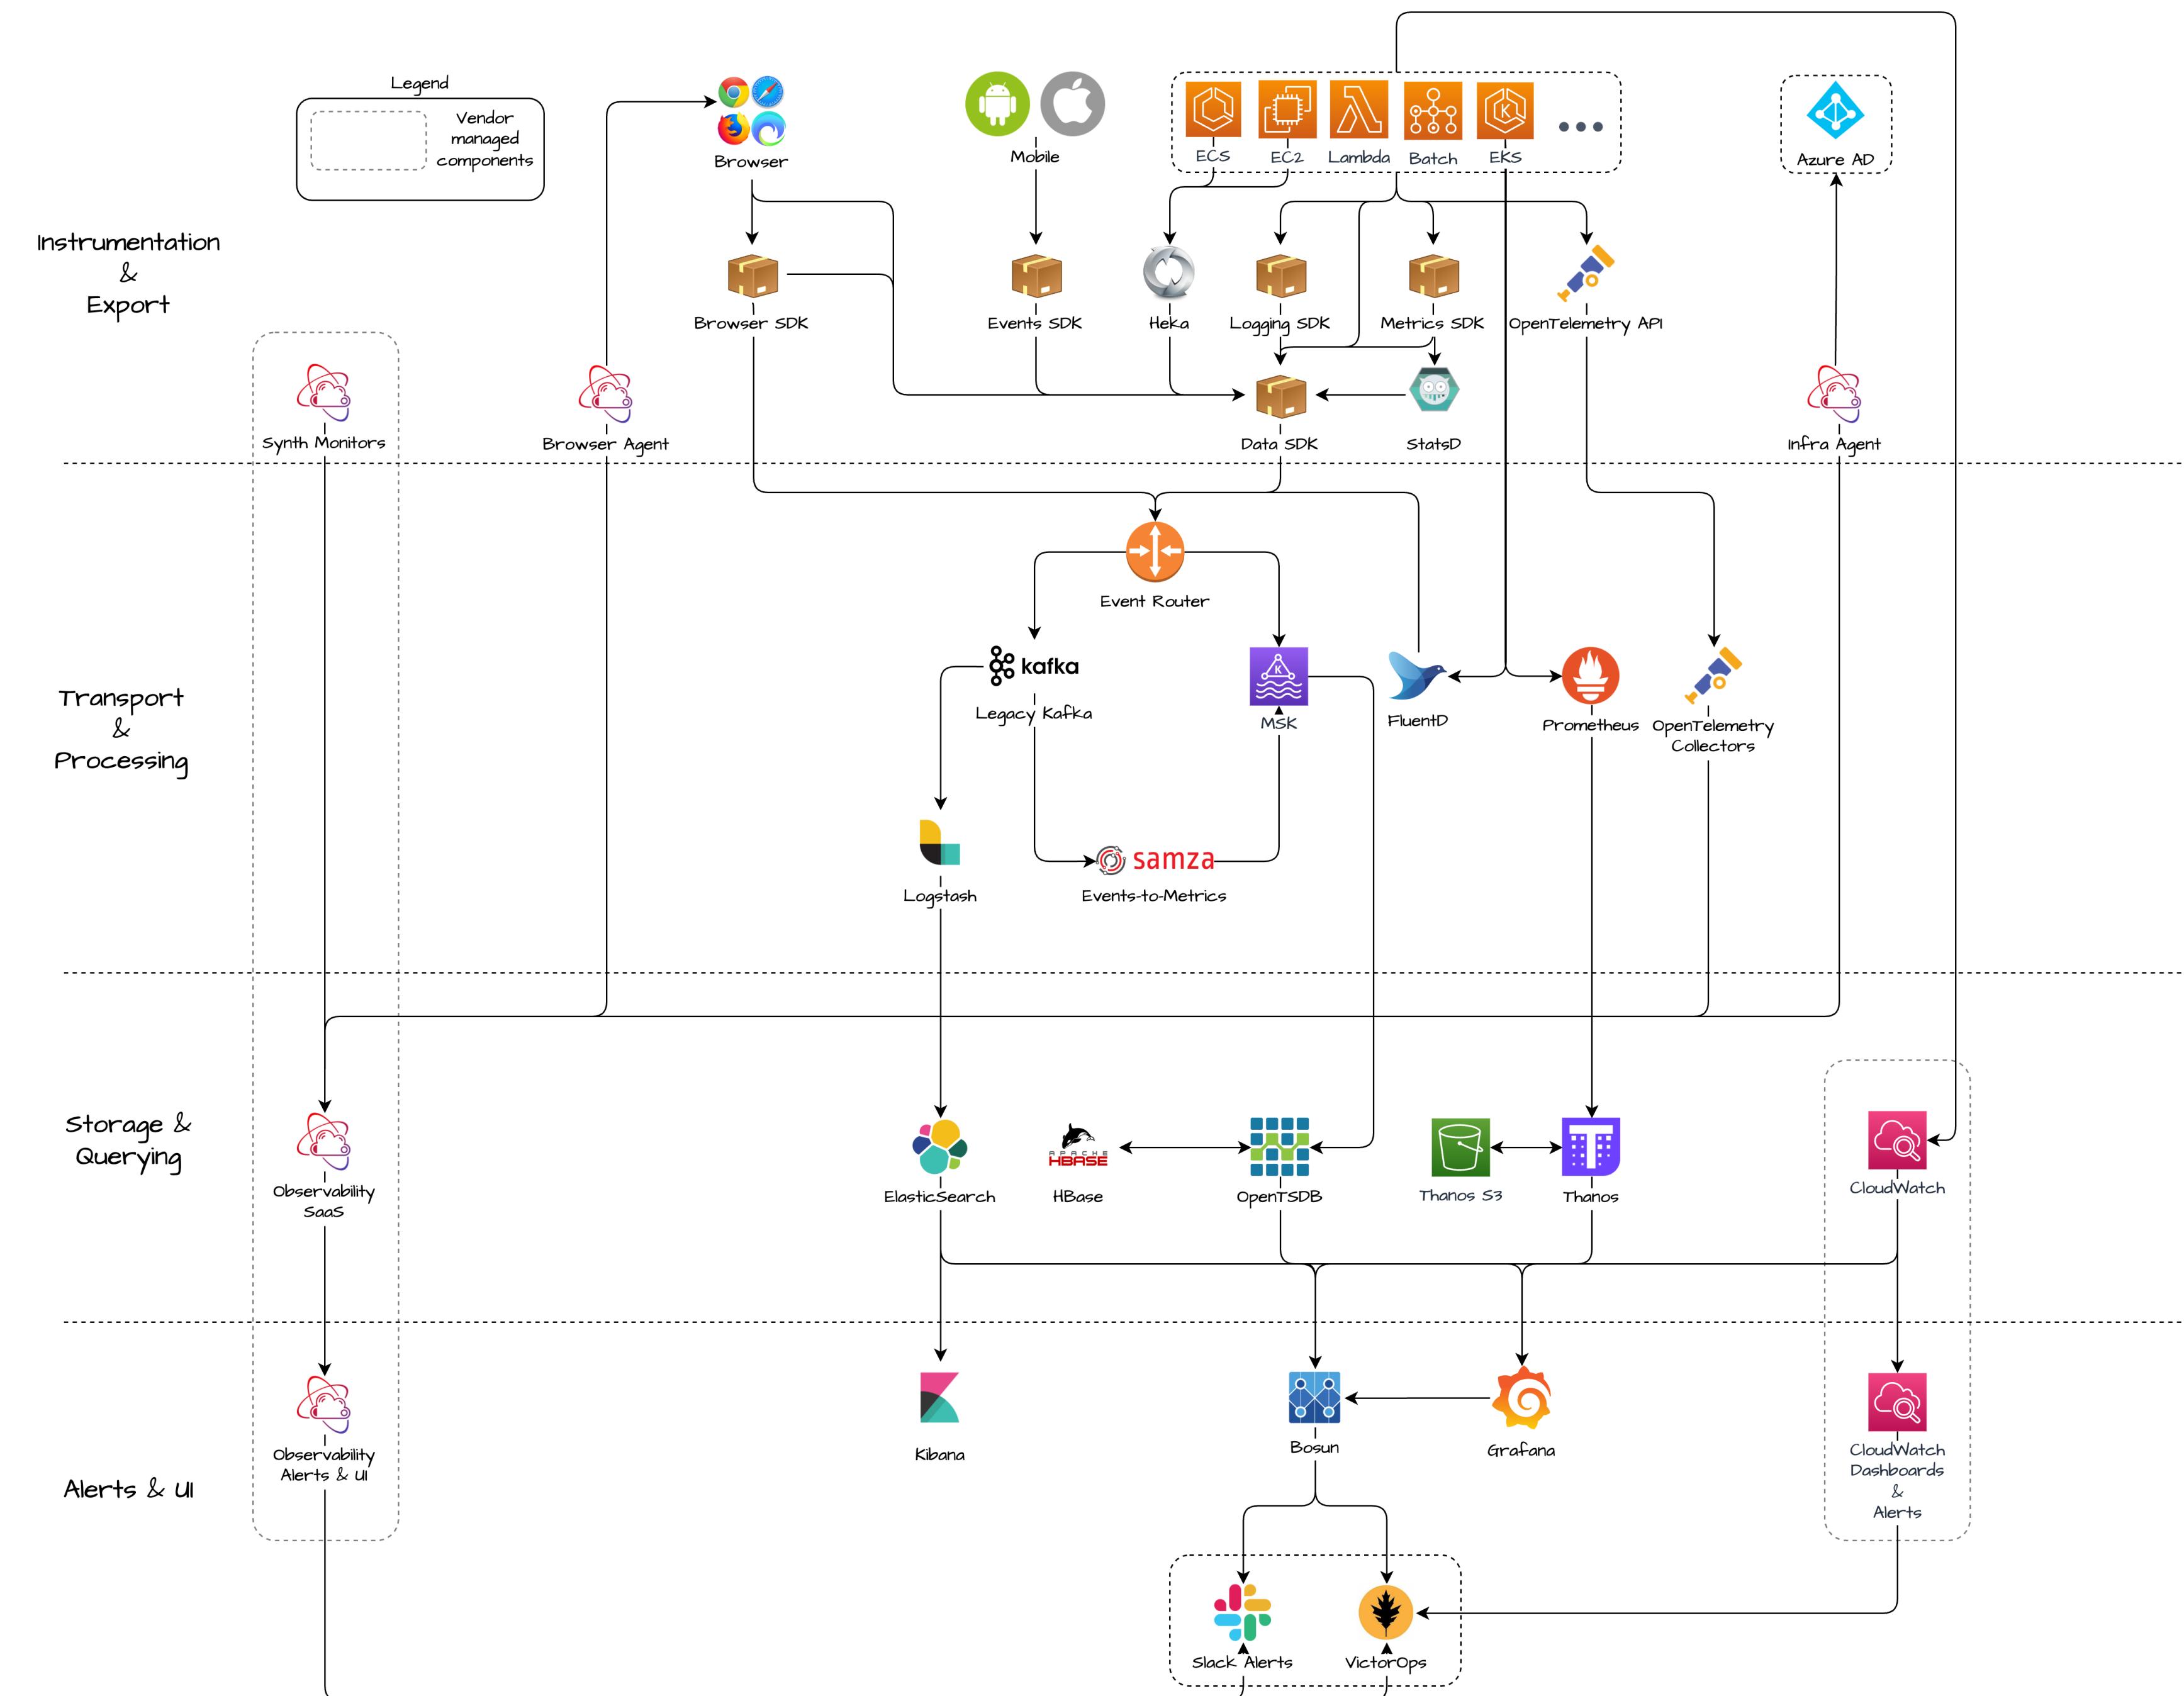
1. Why observability matters
2. How open standards help observability
- 3. Rolling out OpenTelemetry**
4. Adopting observability in practice

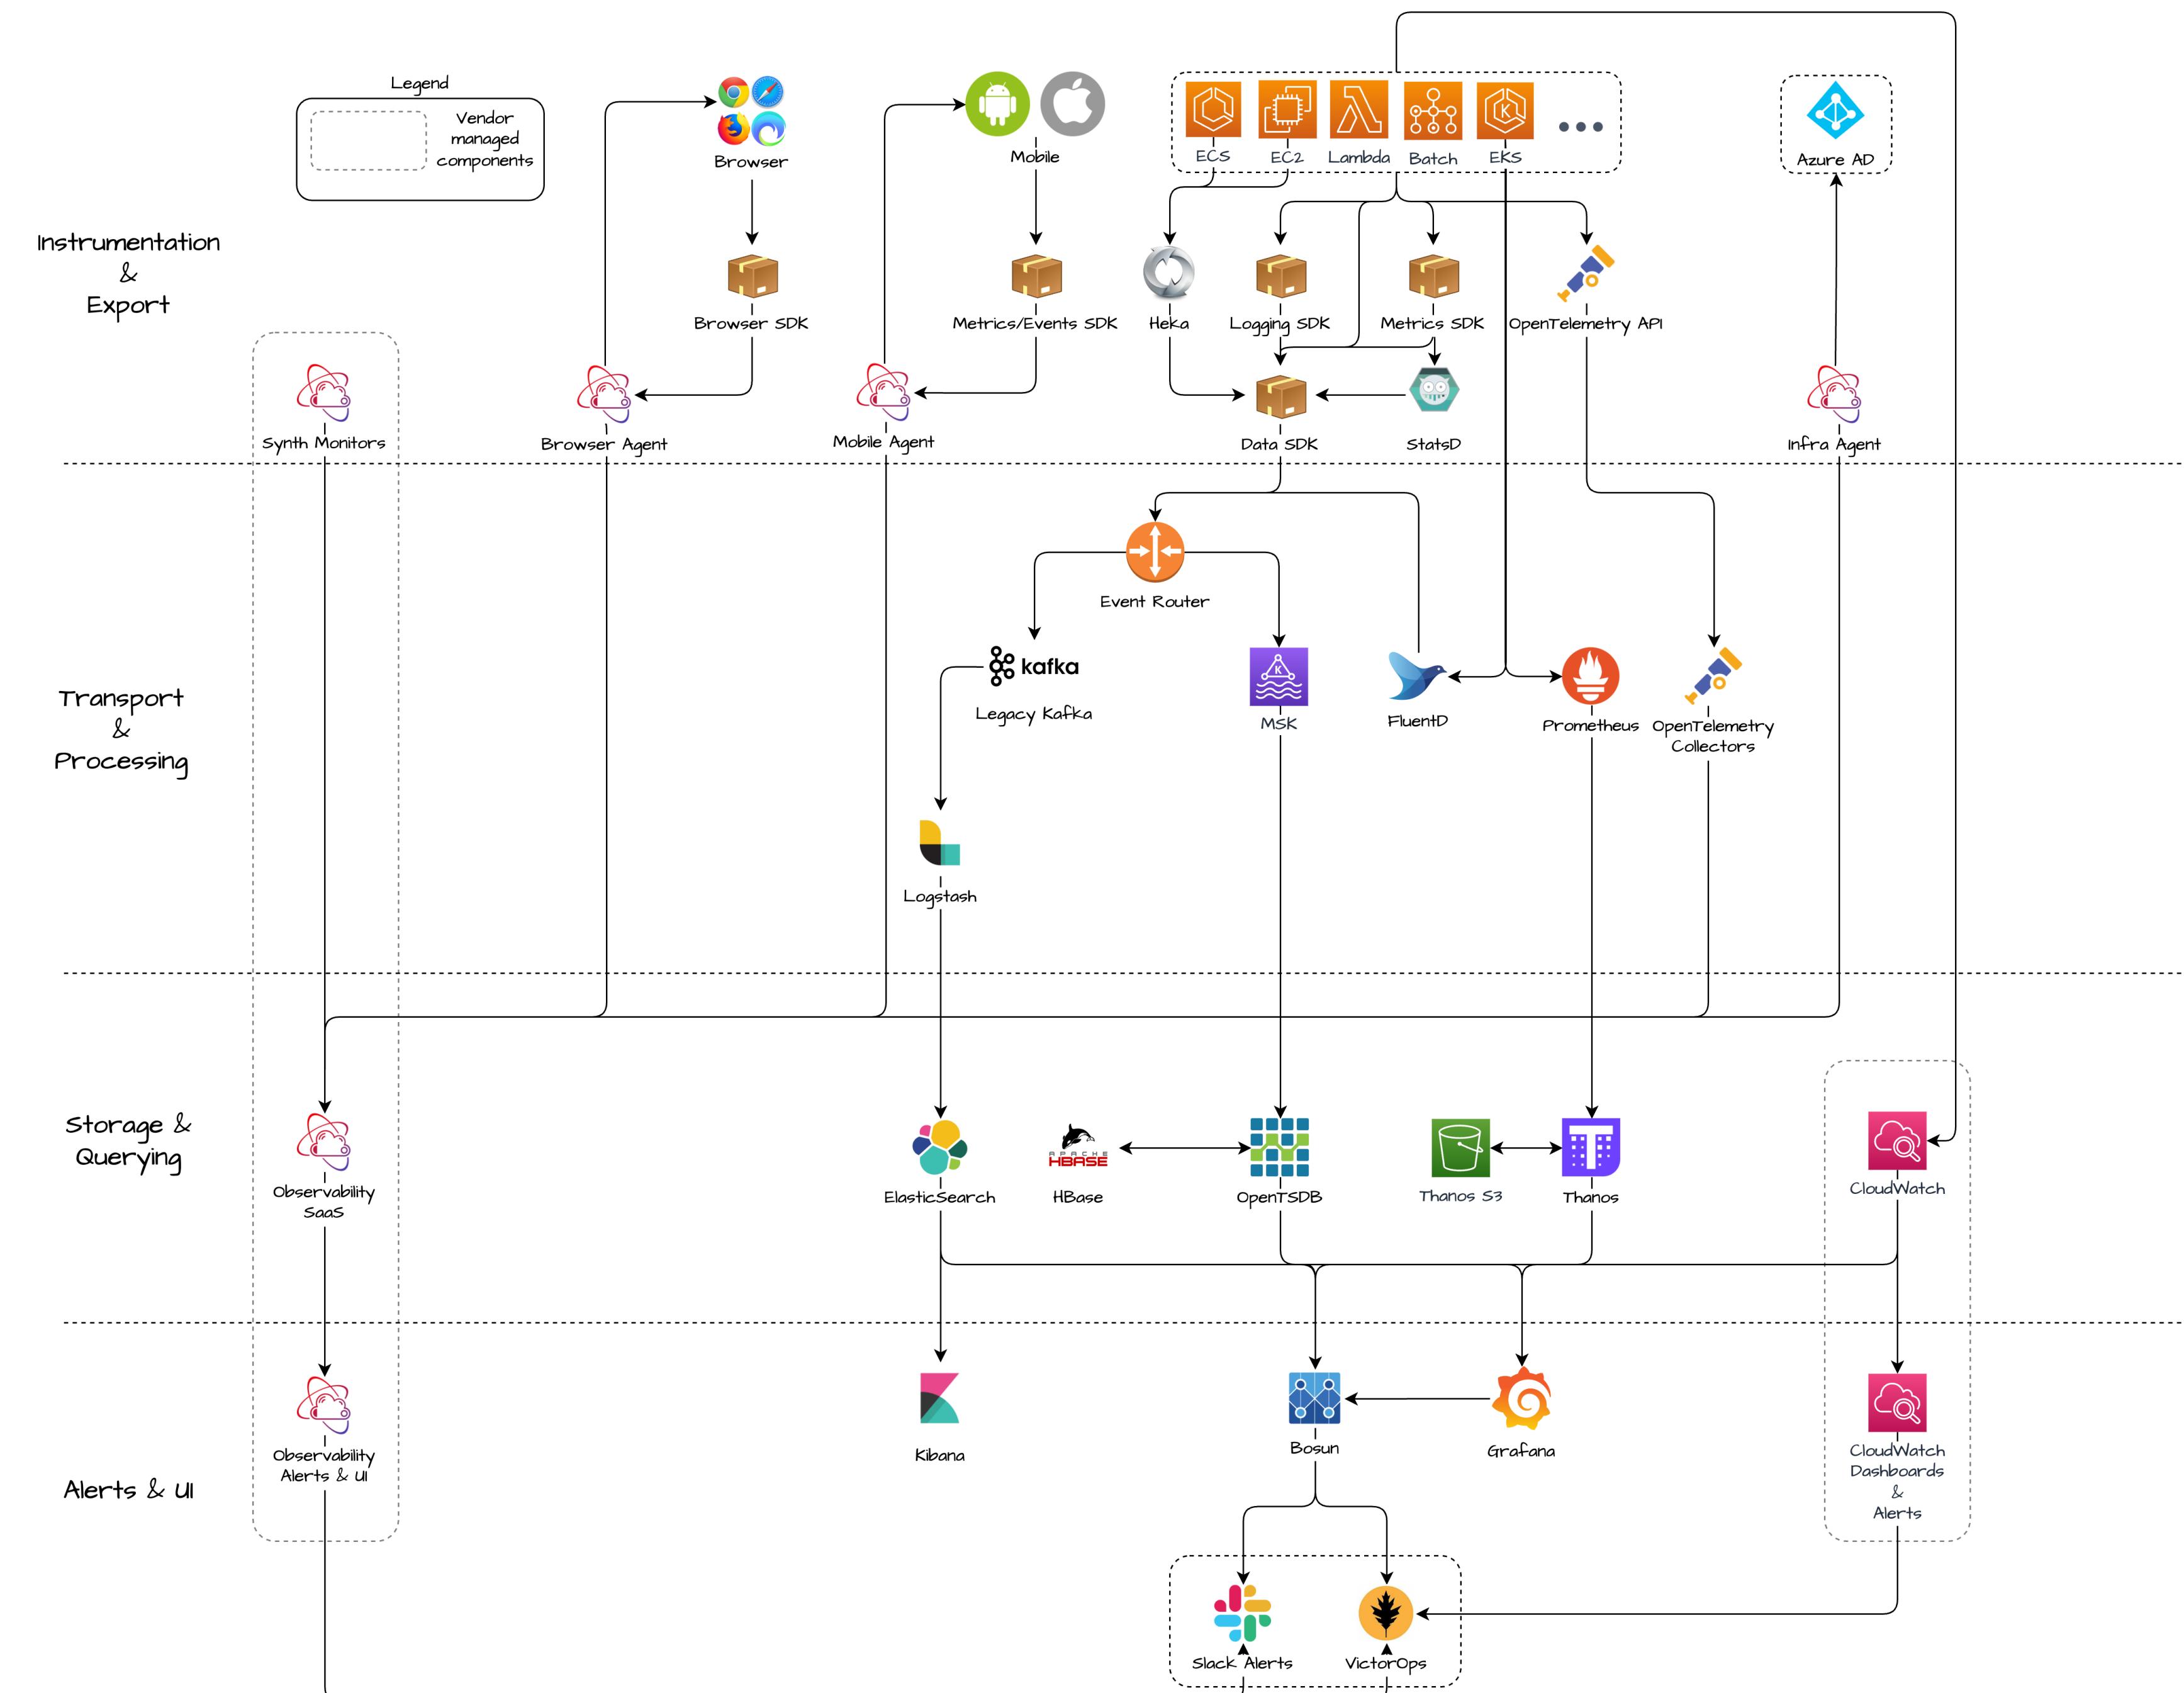
Make the
golden path
the path of
least resistance

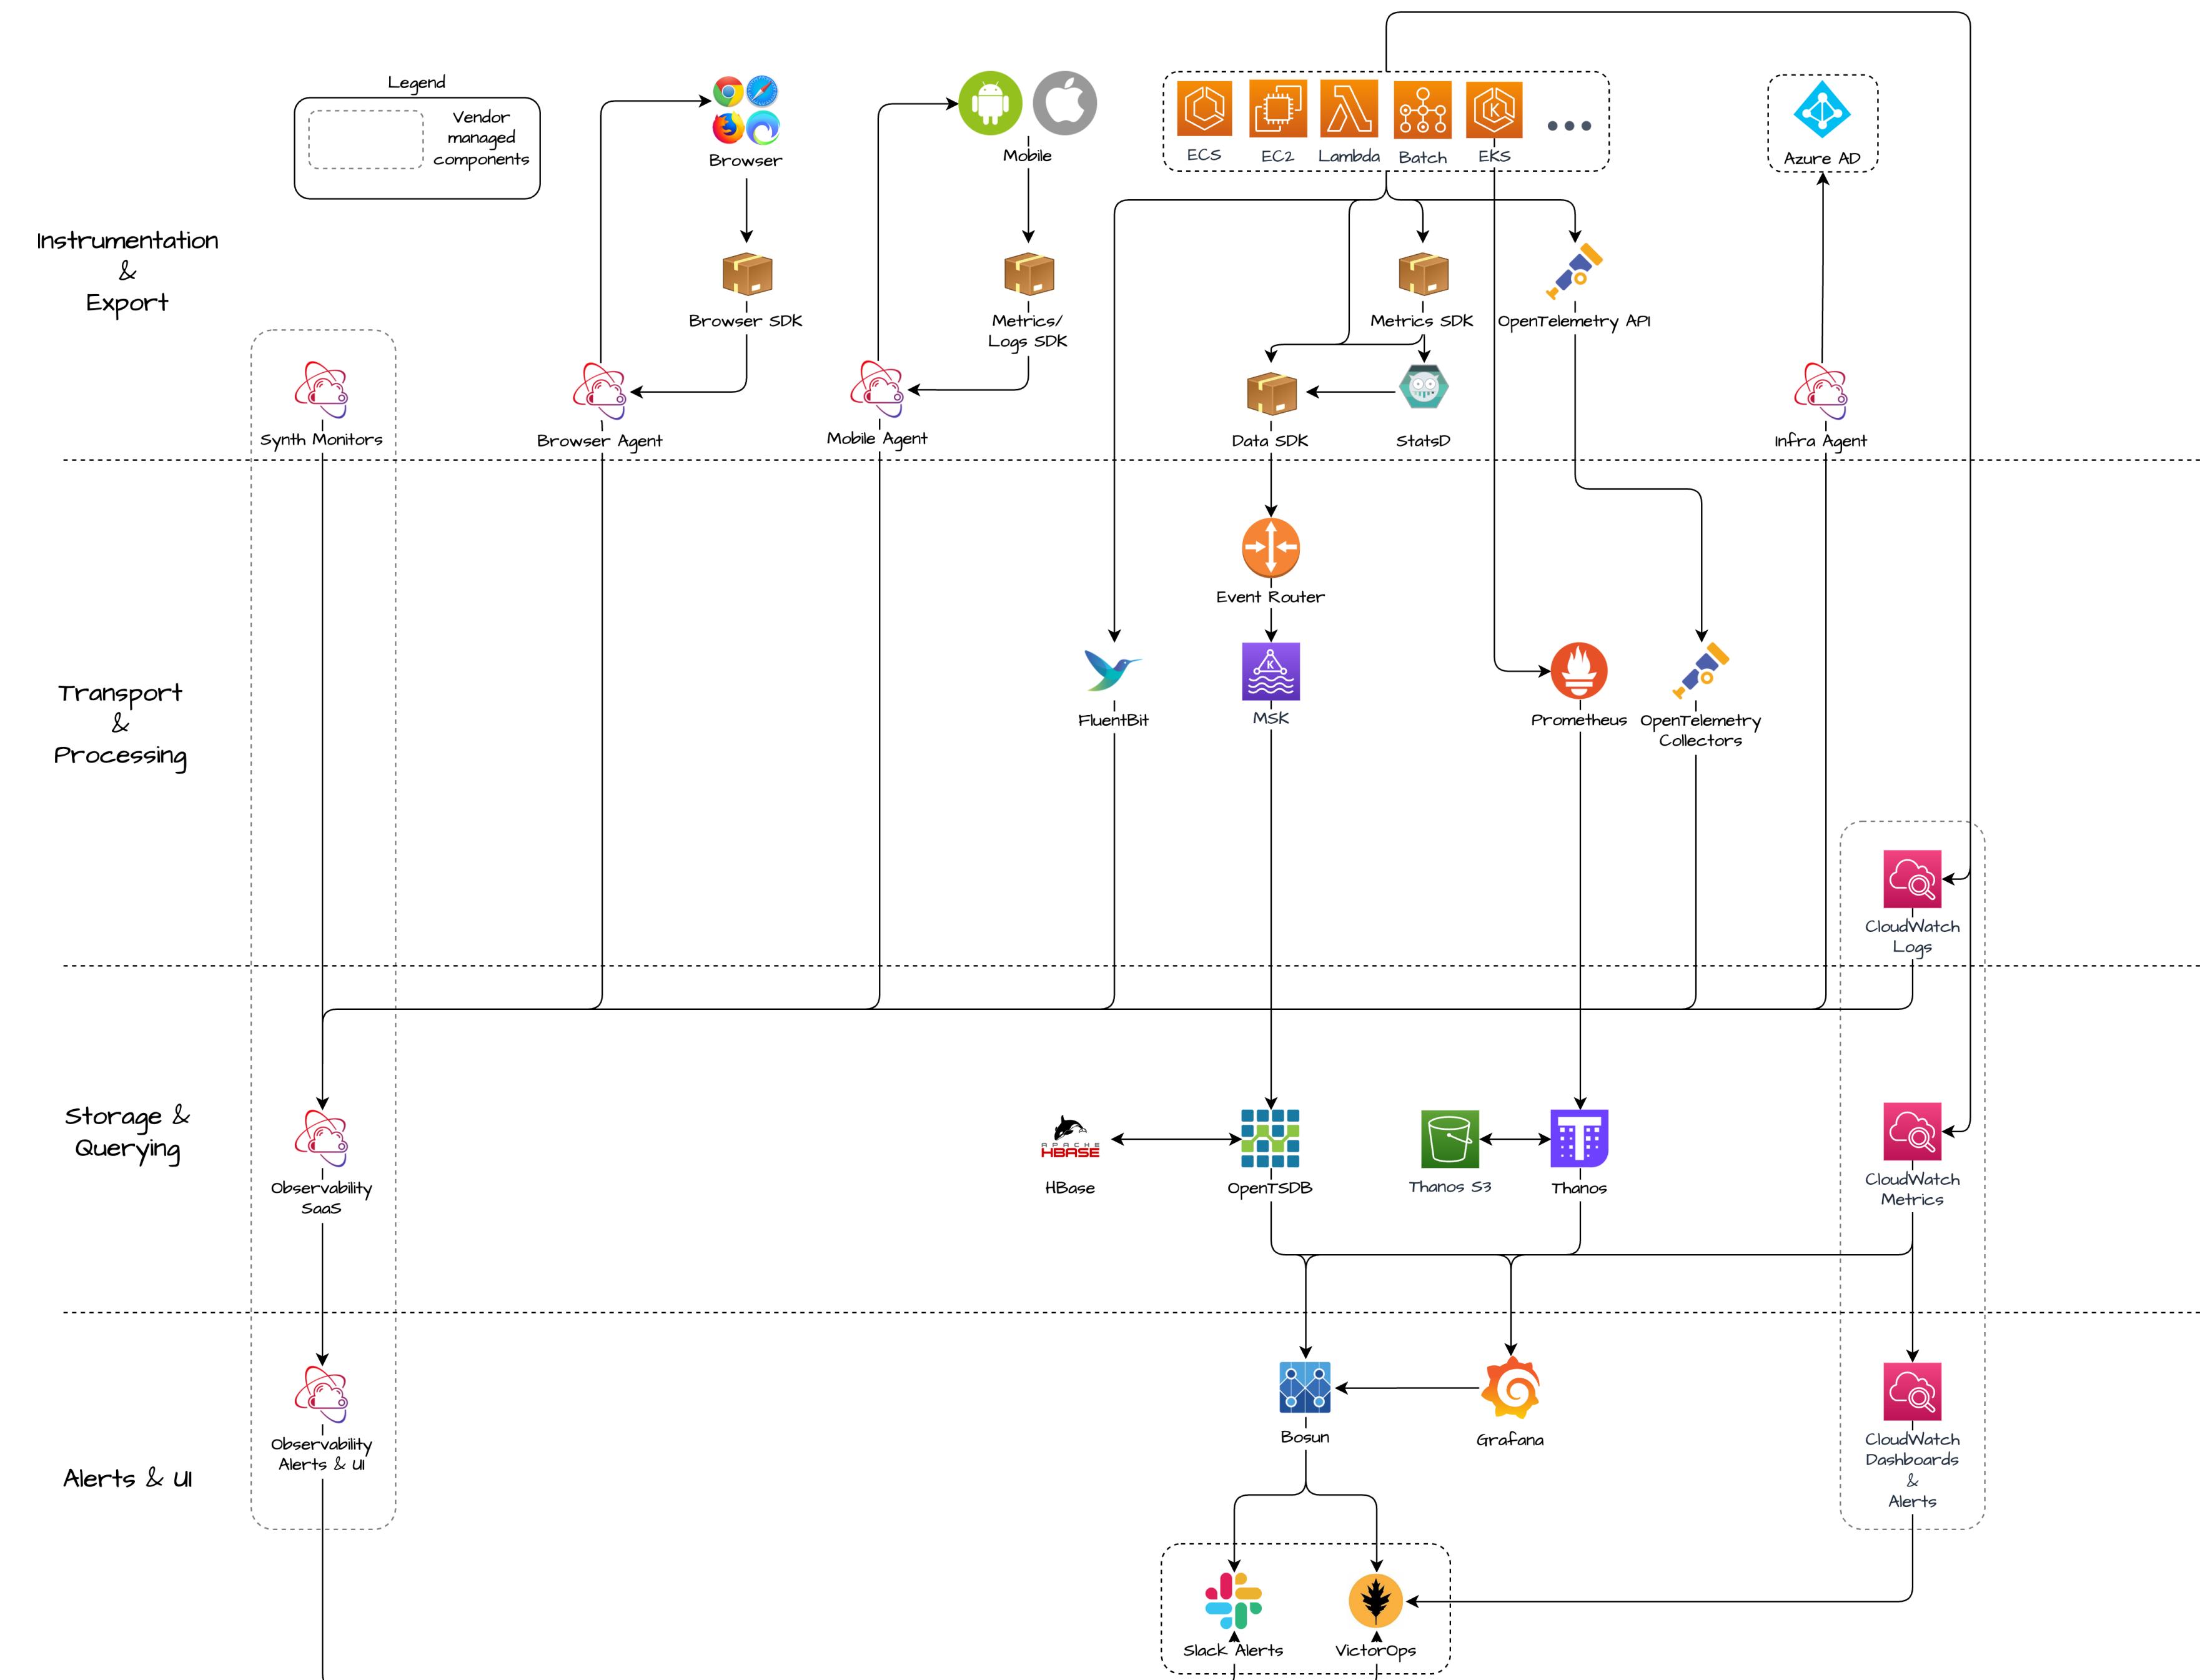


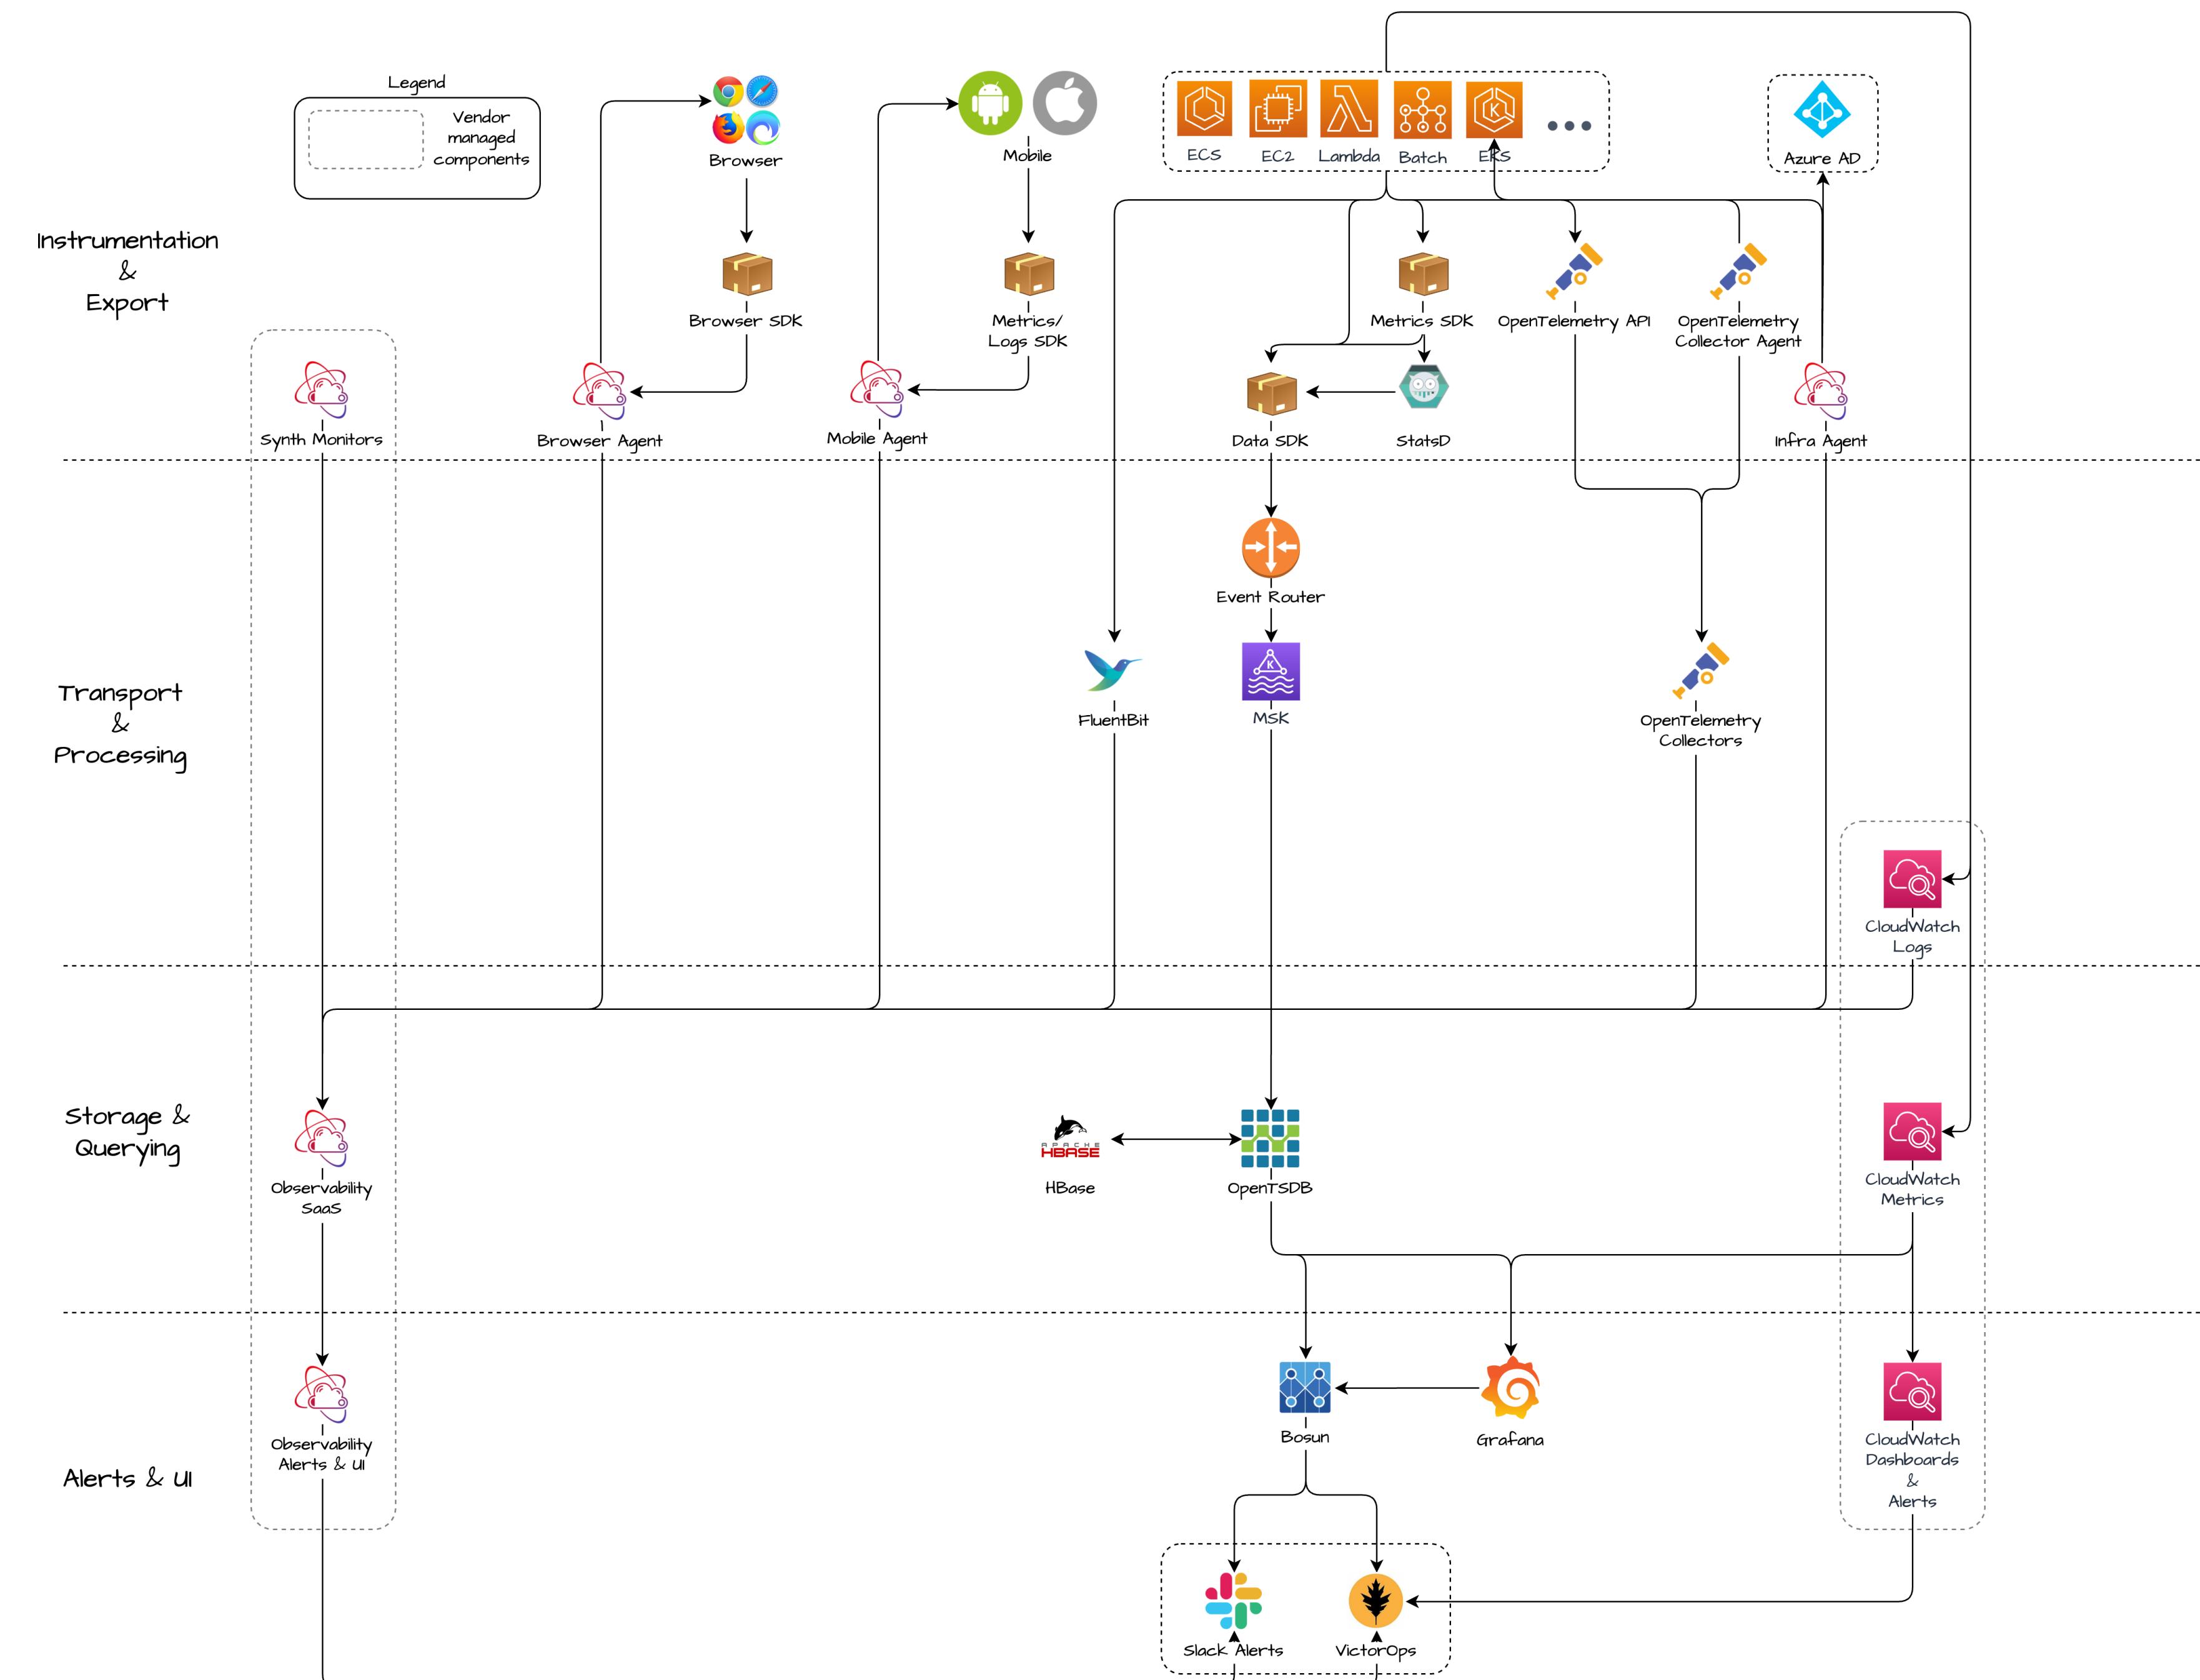


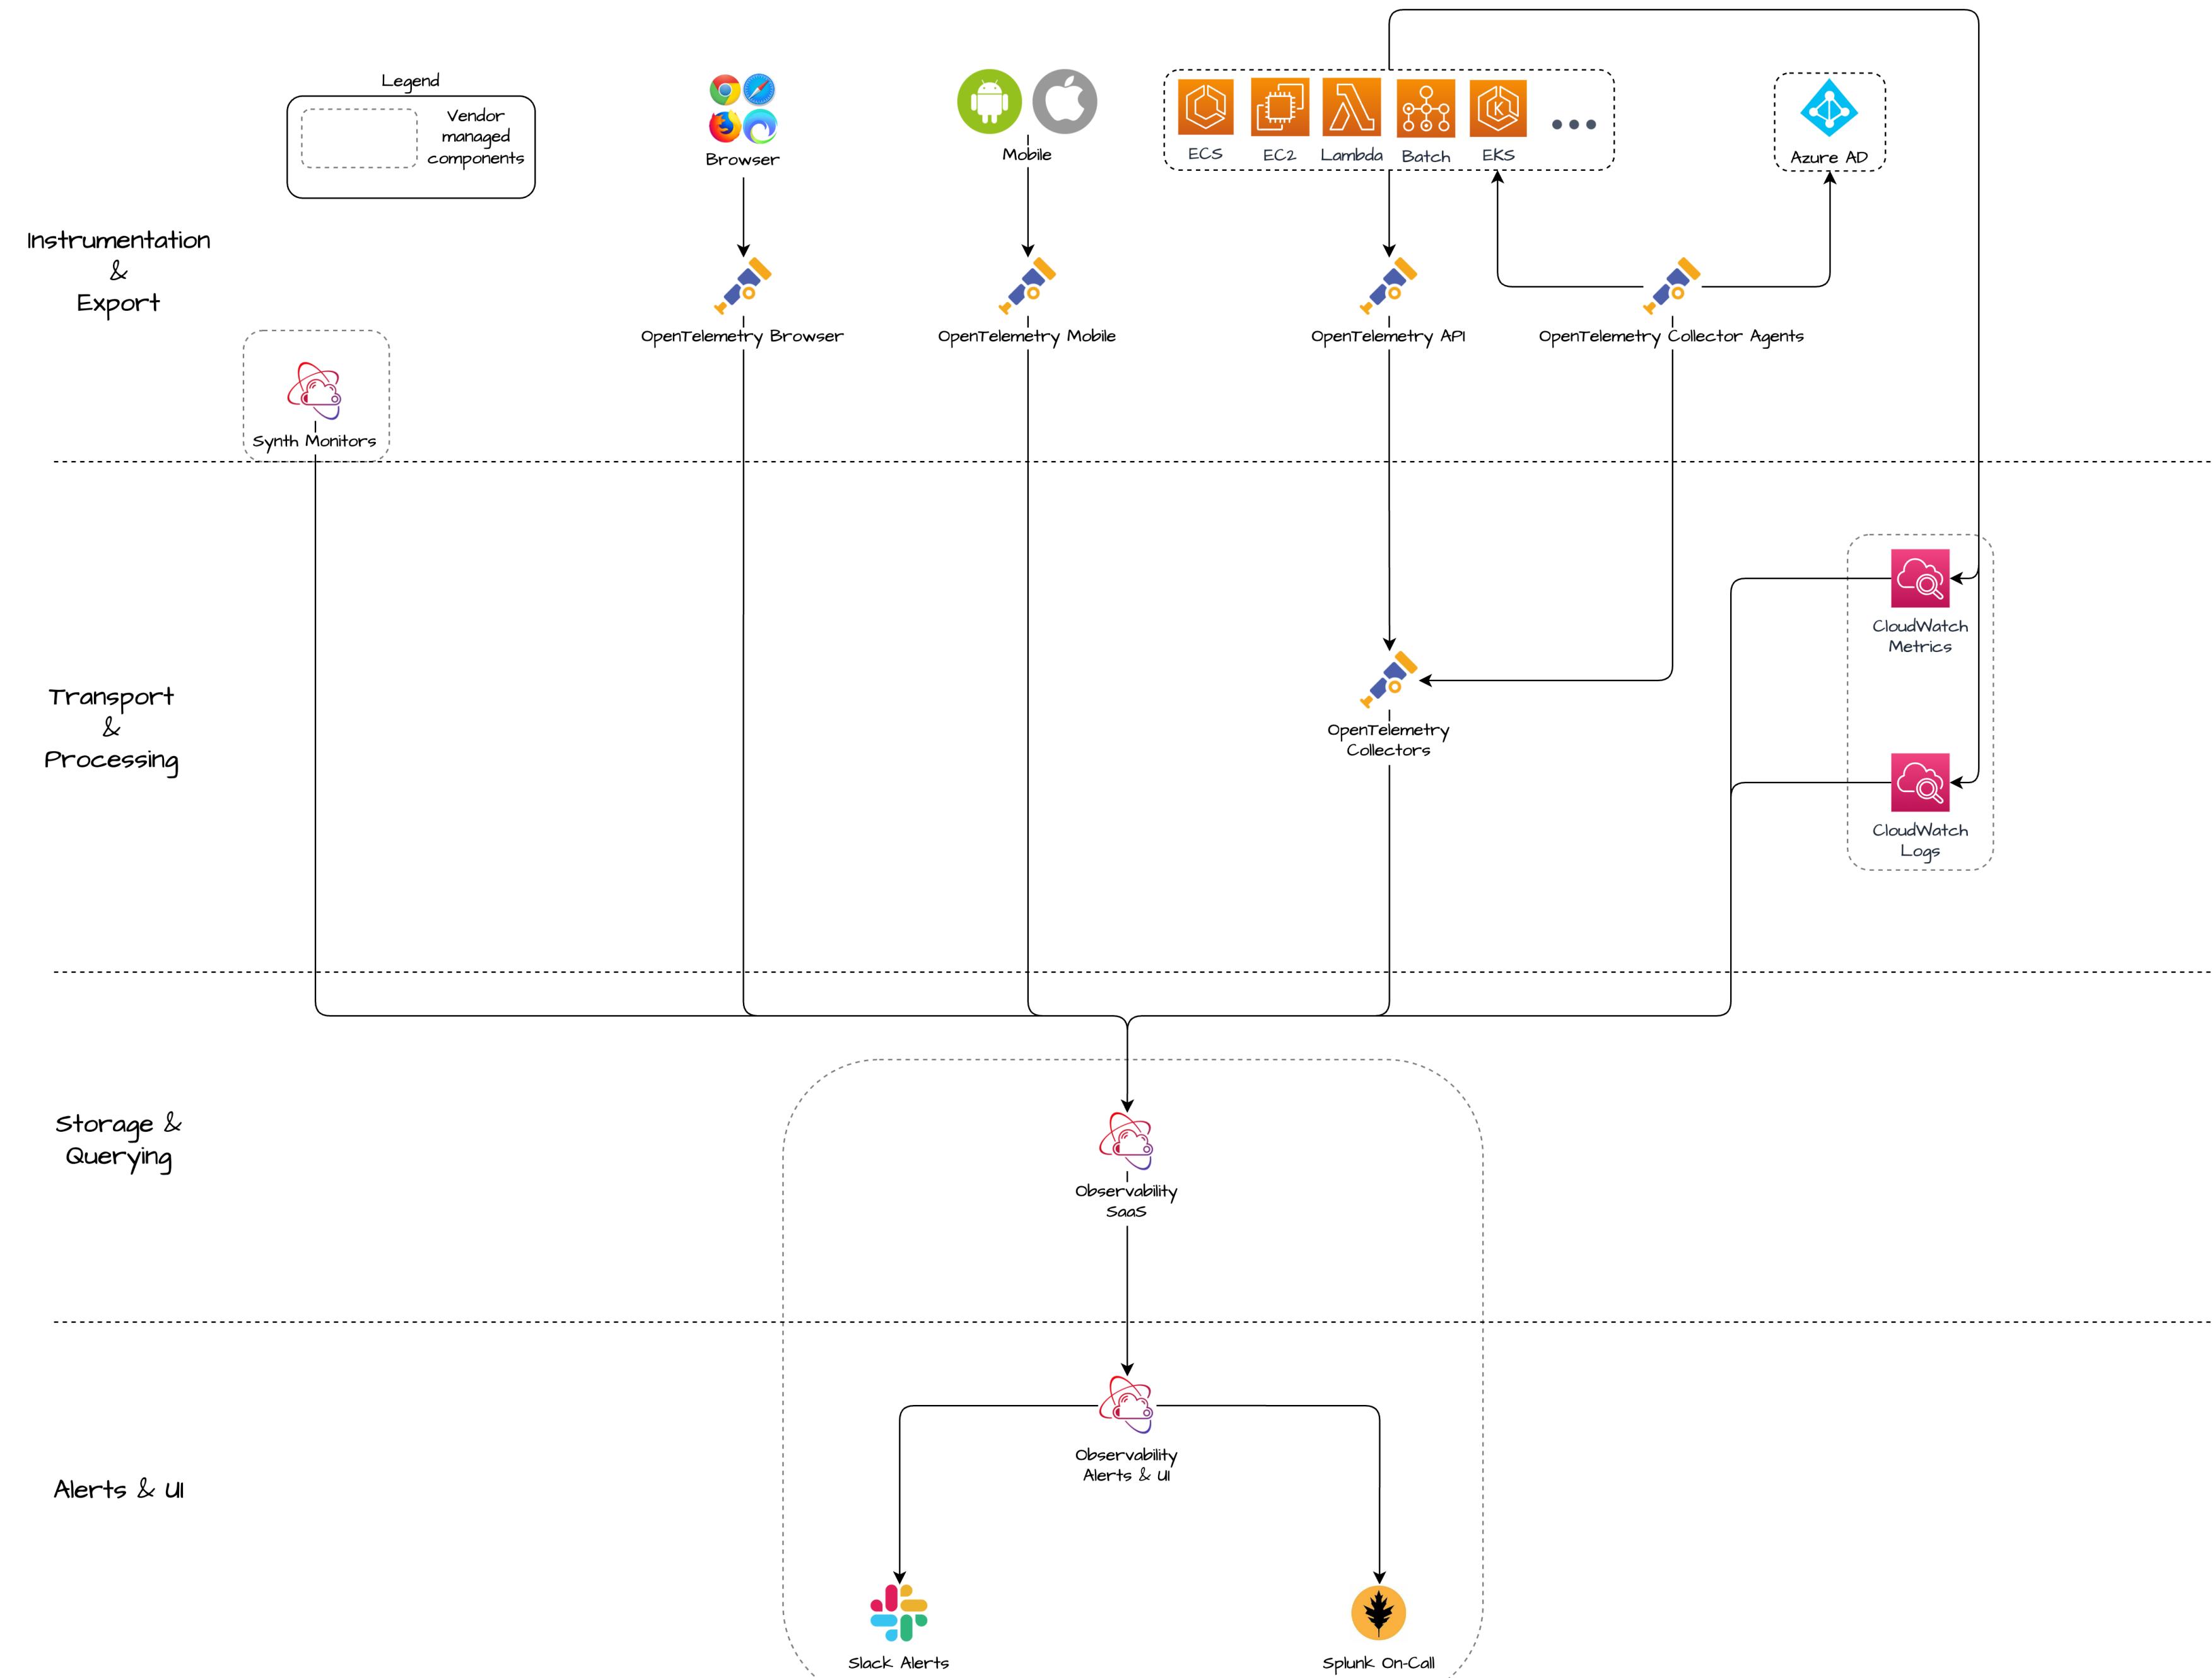




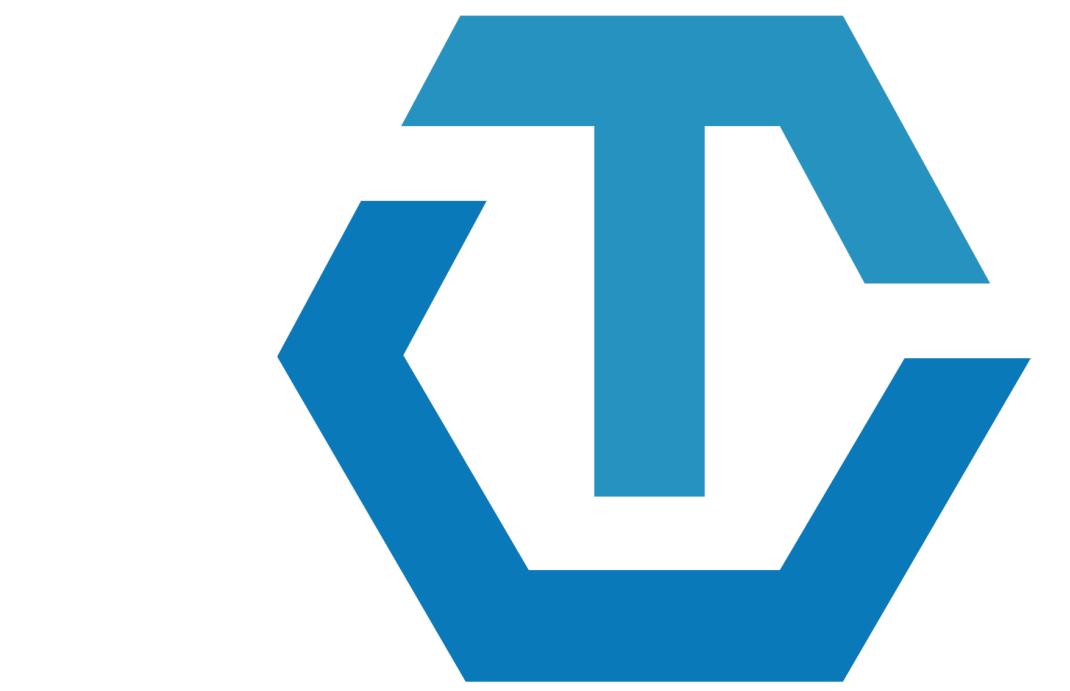




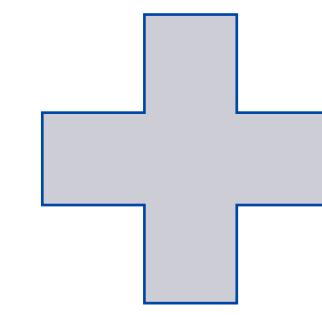




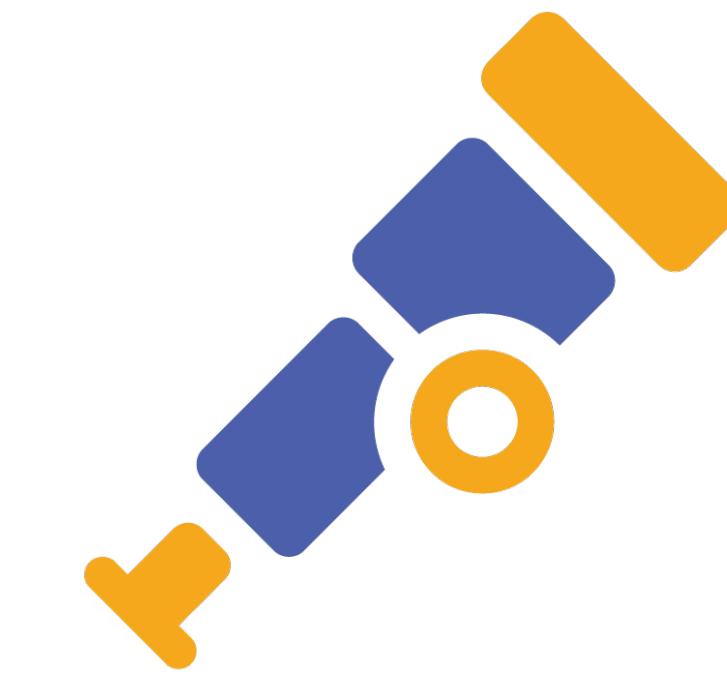
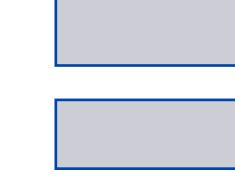
Compatibility with existing solutions



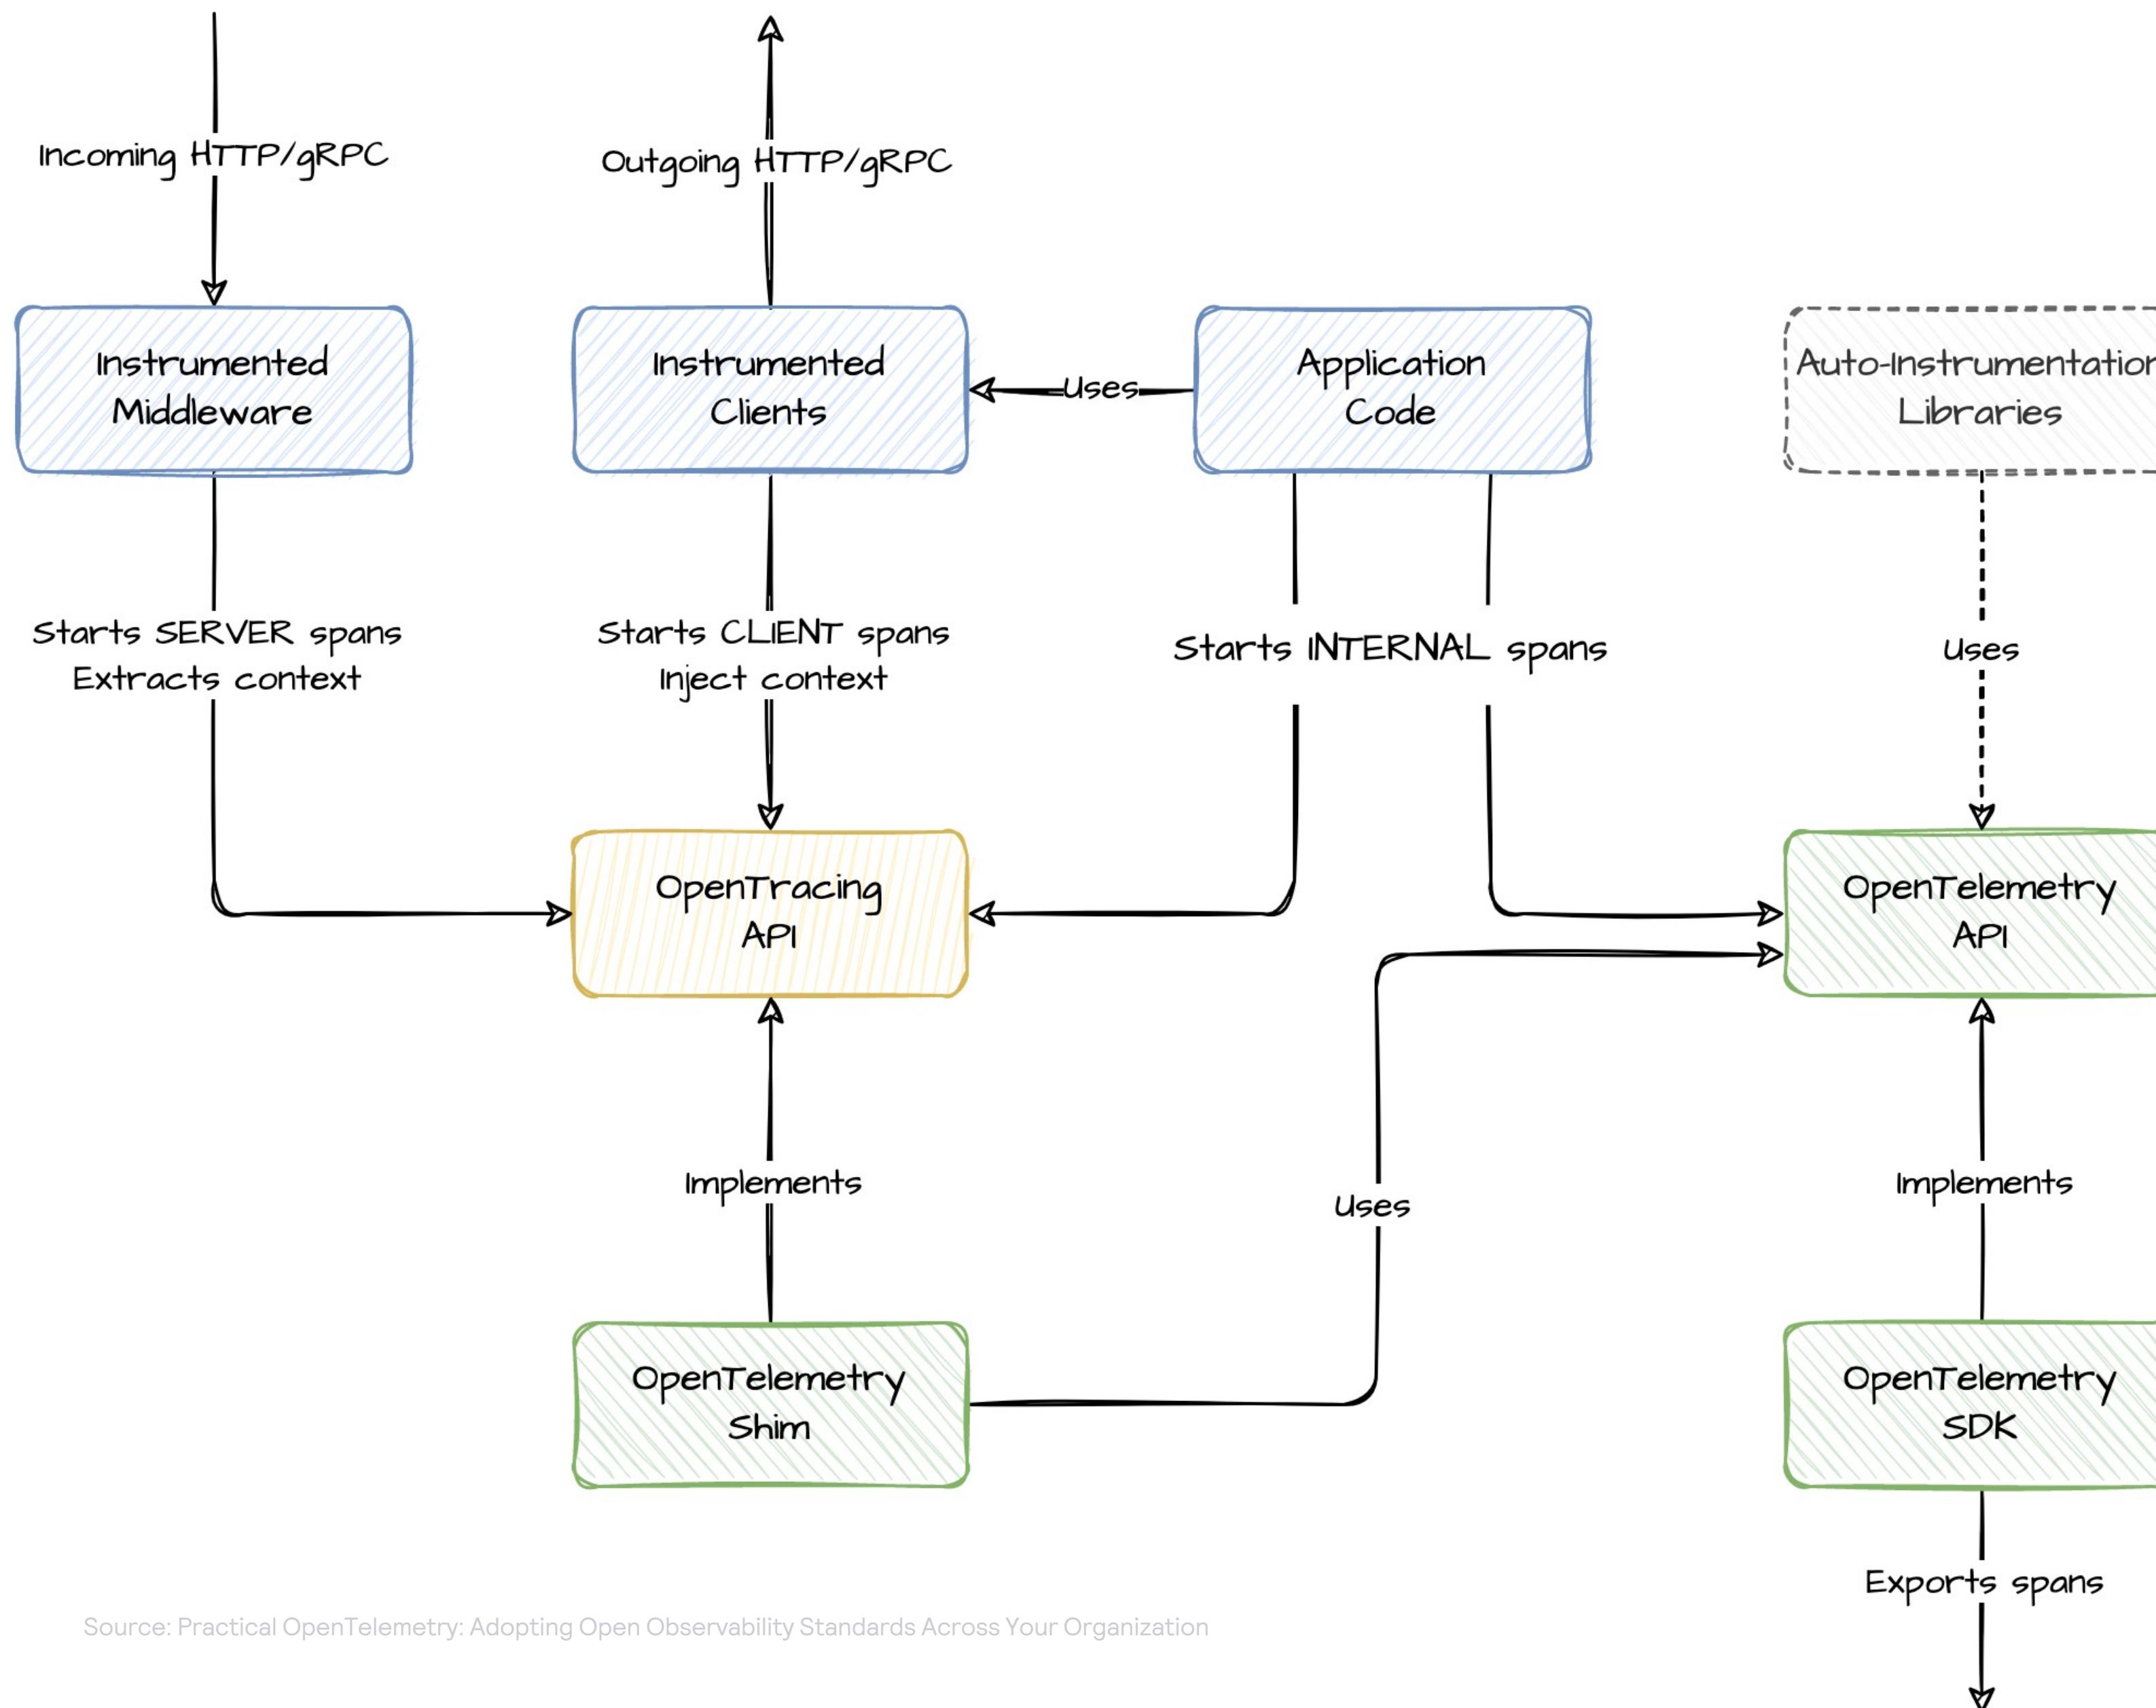
OPEN TRACING



OpenCensus



OpenTelemetry



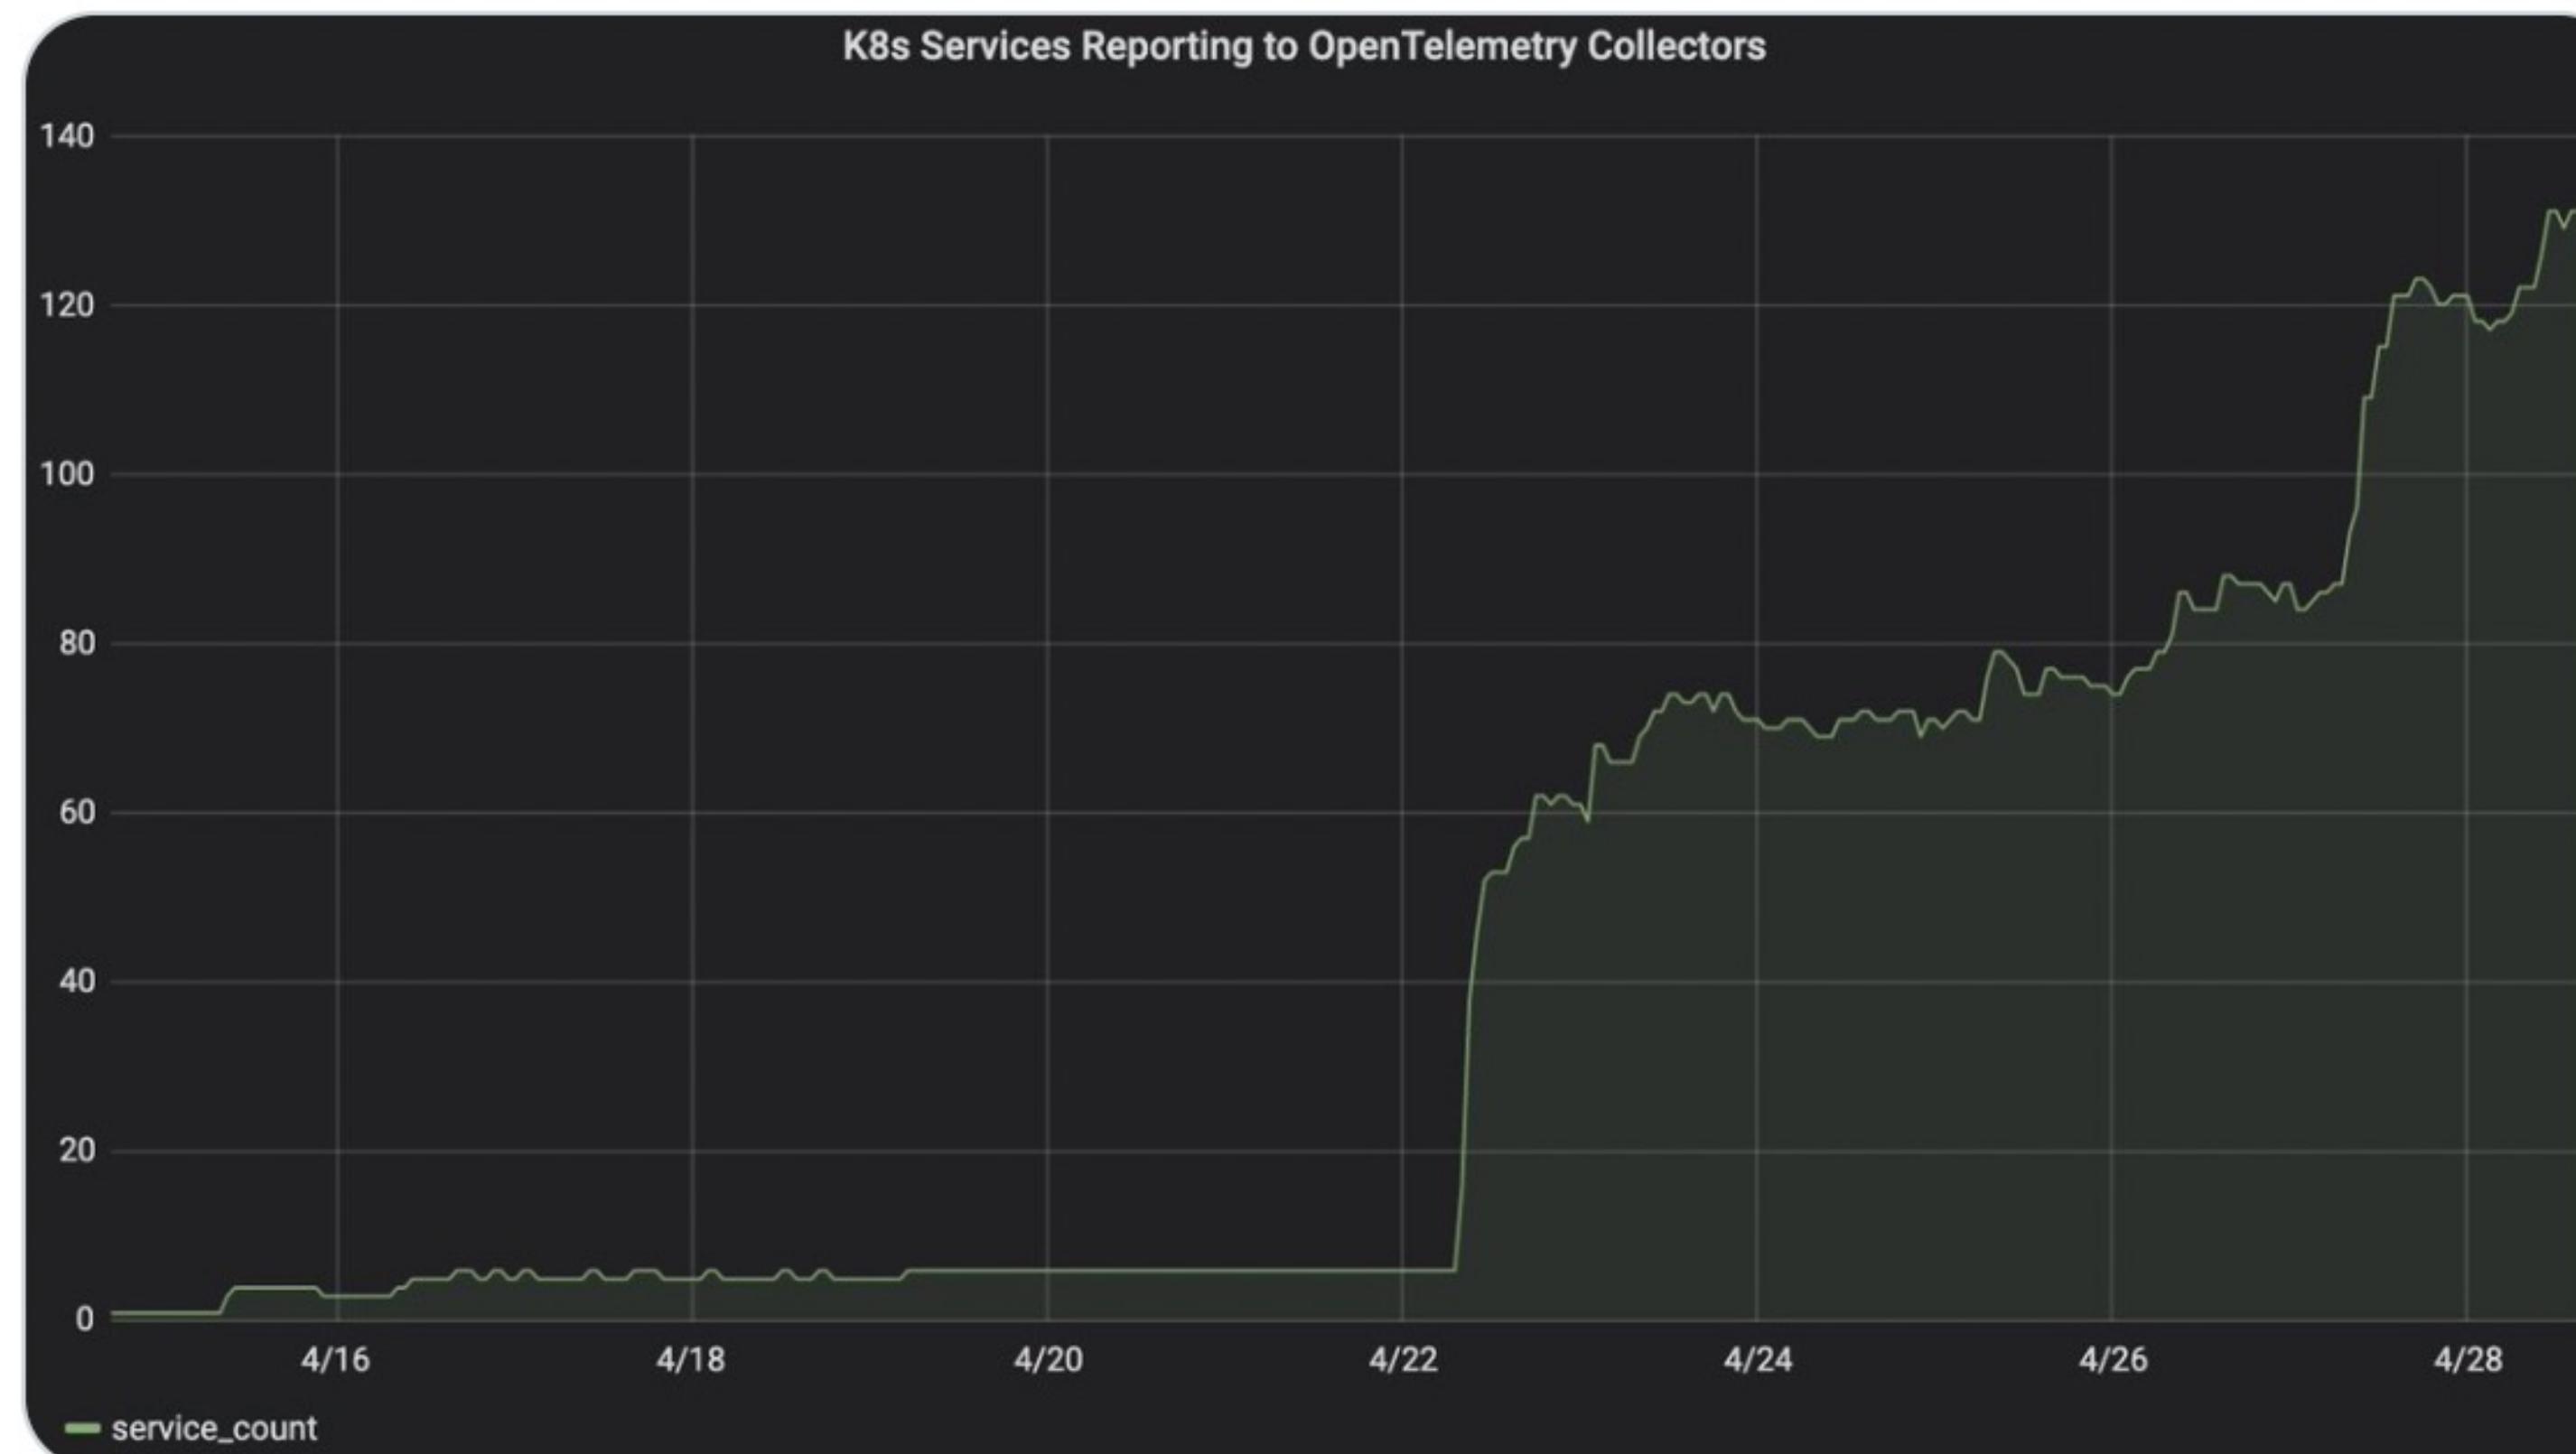
Source: Practical OpenTelemetry: Adopting Open Observability Standards Across Your Organization



Dan Gomez Blanco
@dan_gomezblanco

...

It pays off when your migration to [@opentelemetry](#) involves a minor version bump 😊



9:19 PM · Apr 28, 2021

OTel Collectors are incredibly powerful

Receiving OTLP, Zipkin,
Prometheus

Removing unwanted attributes

Unsetting span status for spans
matching regex

Generate metrics from spans
(bye Istio Mixer)

Converting from cumulative to
delta temporality

Renaming attributes to follow
semantic conventions

Exporting data OTLP and
Prometheus

>1.8M

Spans per second

>90k

Traces per second

<125

Total used CPU cores

<100

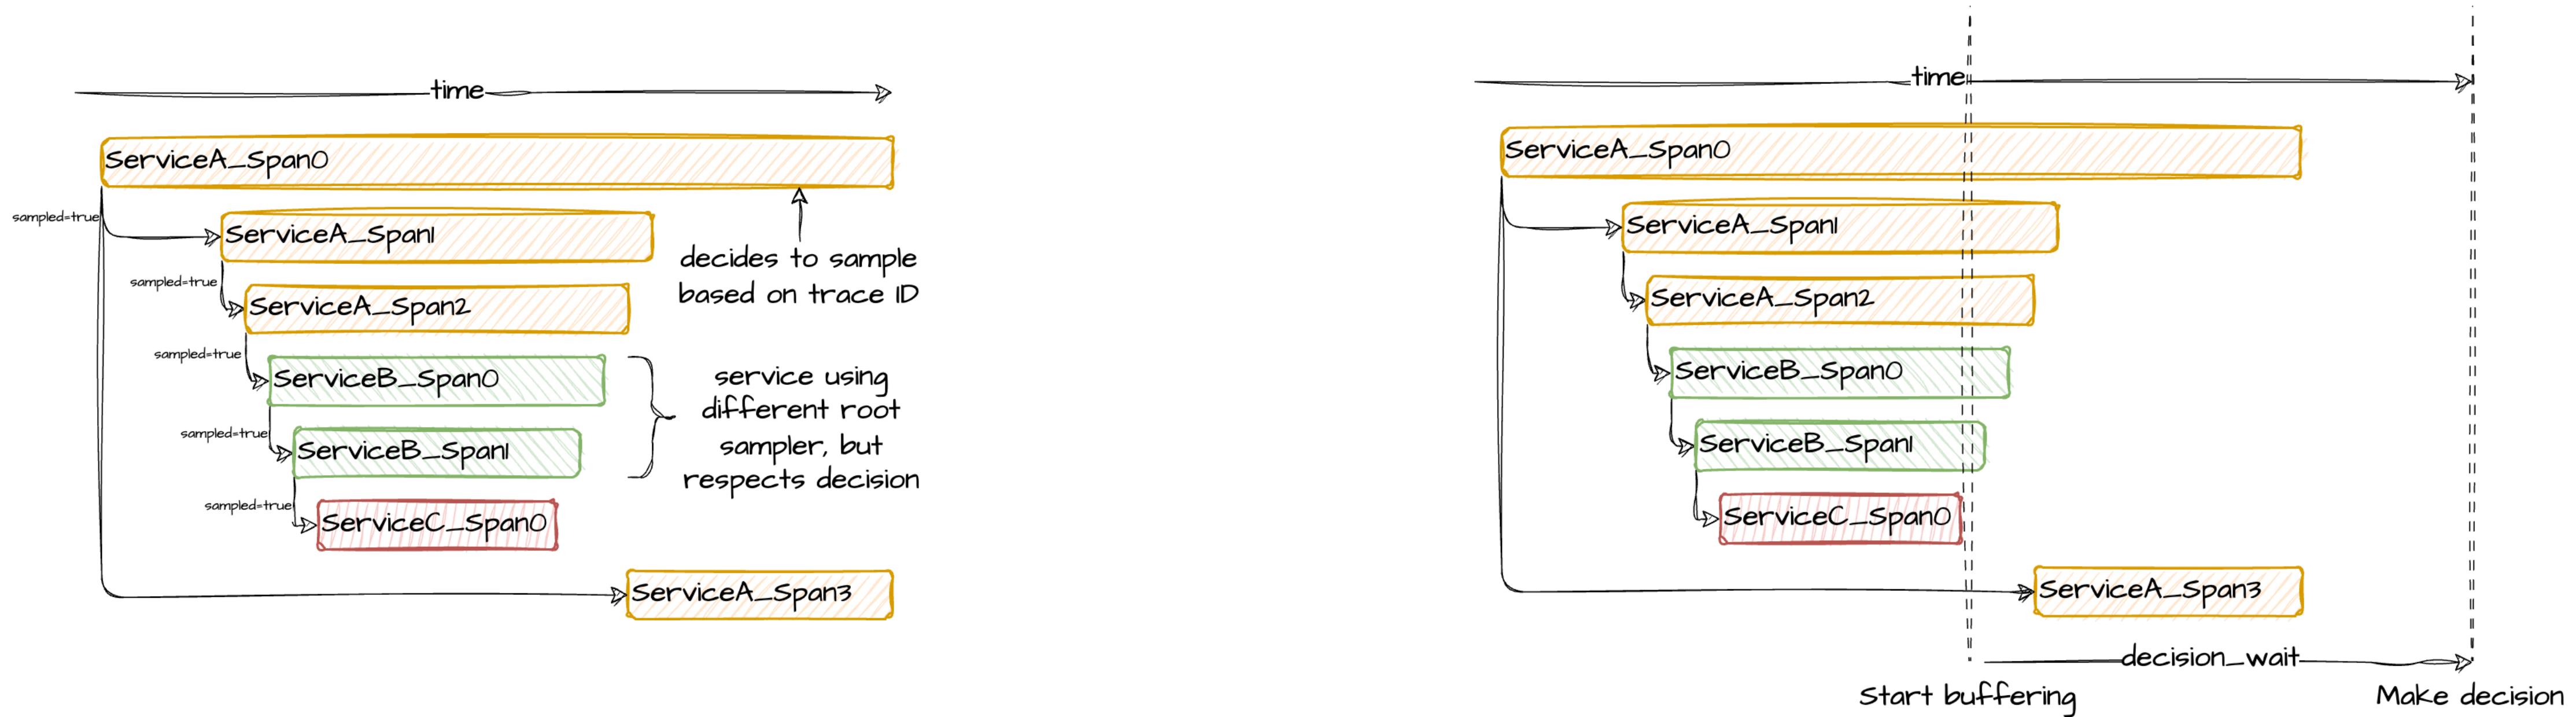
Replicas

Agenda

- 1. Why observability matters**
- 2. How open standards help observability**
- 3. Rolling out OpenTelemetry**
- 4. Adopting observability in practice**

**Most data
gathered for
debug purposes is
never used**

Keep useful data, discard the rest



Probability sampling

A span is sampled based on its properties or the propagated trace context. Simpler to configure.

Tail based sampling

A span is sampled based on properties of the whole trace (e.g. slow traces or those containing errors). More powerful but requires external components.

How much tracing data do we keep?

45%

sampled from all traces

Keeping telemetry data valuable

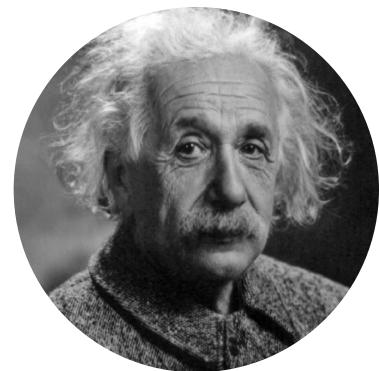
Limiting telemetry production with excessive controls

- Slows down team velocity
- Generates toil for telemetry admins
- Falls out of date soon, defeating its purpose

Visualising cost and encouraging good use of telemetry signals

- Leverage semantic conventions to segment costs to services and namespaces
- Review telemetry along with other service costs
- Reward learning and product health

Failure is success in progress



Albert Einstein

Progress towards success requires action

Foster a learning and improvement culture by:

1. Establishing targets for time-to-detect and time-to-resolve
2. Discussing post-mortems to find areas of improvement
3. Encouraging observability champions to join those discussions
4. Creating a guild/chapter across the company to gather external feedback
5. Running sessions where teams can evaluate telemetry together

- Complex systems require effective observability
- Open standards empower simplification
- OpenTelemetry enables signals to be used efficiently

Thank you